THE RUBY ON RAILS TUTORIAL

LEARN WEB DEVELOPMENT WITH RAILS

THIRD EDITION

BOOK AND SCREENCASTS BY MICHAEL HARTL

2

Ruby on Rails Tutorial

Learn Web Development with Rails

Michael Hartl

ii

Contents

Foreword

My former company (CD Baby) was one of the first to loudly switch to Ruby on Rails, and then even more loudly switch back to PHP (Google me to read about the drama). This book by Michael Hartl came so highly recommended that I had to try it, and the *Ruby on Rails Tutorial* is what I used to switch back to Rails again.

Though I've worked my way through many Rails books, this is the one that finally made me "get" it. Everything is done very much "the Rails way"—a way that felt very unnatural to me before, but now after doing this book finally feels natural. This is also the only Rails book that does test-driven development the entire time, an approach highly recommended by the experts but which has never been so clearly demonstrated before. Finally, by including Git, GitHub, and Heroku in the demo examples, the author really gives you a feel for what it's like to do a real-world project. The tutorial's code examples are not in isolation.

The linear narrative is such a great format. Personally, I powered through the *Rails Tutorial* in three long days,[1] doing all the examples and challenges at the end of each chapter. Do it from start to finish, without jumping around, and you'll get the ultimate benefit.

Enjoy!

Derek Sivers (sivers.org)
Founder, CD Baby

---

[1]This is not typical! Getting through the entire book usually takes *much* longer than three days.

Acknowledgments

The *Ruby on Rails Tutorial* owes a lot to my previous Rails book, *RailsSpace*, and hence to my coauthor Aurelius Prochazka. I'd like to thank Aure both for the work he did on that book and for his support of this one. I'd also like to thank Debra Williams Cauley, my editor on both *RailsSpace* and the *Ruby on Rails Tutorial*; as long as she keeps taking me to baseball games, I'll keep writing books for her.

I'd like to acknowledge a long list of Rubyists who have taught and inspired me over the years: David Heinemeier Hansson, Yehuda Katz, Carl Lerche, Jeremy Kemper, Xavier Noria, Ryan Bates, Geoffrey Grosenbach, Peter Cooper, Matt Aimonetti, Mark Bates, Gregg Pollack, Wayne E. Seguin, Amy Hoy, Dave Chelimsky, Pat Maddox, Tom Preston-Werner, Chris Wanstrath, Chad Fowler, Josh Susser, Obie Fernandez, Ian McFarland, Steven Bristol, Pratik Naik, Sarah Mei, Sarah Allen, Wolfram Arnold, Alex Chaffee, Giles Bowkett, Evan Dorn, Long Nguyen, James Lindenbaum, Adam Wiggins, Tikhon Bernstam, Ron Evans, Wyatt Greene, Miles Forrest, Sandi Metz, Ryan Davis, Aaron Patterson, the good people at Pivotal Labs, the Heroku gang, the thoughtbot guys, and the GitHub crew. Finally, many, many readers—far too many to list—have contributed a huge number of bug reports and suggestions during the writing of this book, and I gratefully acknowledge their help in making it as good as it can be.

About the author

Michael Hartl is the author of the *Ruby on Rails Tutorial*, one of the leading introductions to web development, and is a cofounder of the Softcover self-publishing platform. His prior experience includes writing and developing *RailsSpace*, an extremely obsolete Rails tutorial book, and developing Insoshi, a once-popular and now-obsolete social networking platform in Ruby on Rails. In 2011, Michael received a Ruby Hero Award for his contributions to the Ruby community. He is a graduate of Harvard College, has a Ph.D. in Physics from Caltech, and is an alumnus of the Y Combinator entrepreneur program.

Copyright and license

Ruby on Rails Tutorial: Learn Web Development with Rails. Copyright © 2014 by Michael Hartl. All source code in the *Ruby on Rails Tutorial* is available jointly under the MIT License and the Beerware License.

Chapter 1

From zero to deploy

Welcome to *Ruby on Rails Tutorial: Learn Web Development with Rails*. The purpose of this book is to teach you how to develop custom web applications, and our tool of choice is the popular Ruby on Rails web framework. If you are new to the subject, the *Ruby on Rails Tutorial* will give you a thorough introduction to web application development, including a basic grounding in Ruby, Rails, HTML & CSS, databases, version control, testing, and deployment—sufficient to launch you on a career as a web developer or technology entrepreneur. On the other hand, if you already know web development, this book will quickly teach you the essentials of the Rails framework, including MVC and REST, generators, migrations, routing, and embedded Ruby. In either case, when you finish the *Ruby on Rails Tutorial* you will be in a position to benefit from the many more advanced books, blogs, and screencasts that are part of the thriving programming educational ecosystem.[1]

The *Ruby on Rails Tutorial* takes an integrated approach to web development by building three example applications of increasing sophistication, starting with a minimal *hello* app (Section 1.3), a slightly more capable *toy* app (Chapter 2), and a real *sample* app (Chapter 3 through Chapter 12). As implied by their generic names, the applications developed in the *Ruby on Rails Tutorial* are not specific to any particular kind of website; although the final

---

[1]The most up-to-date version of the *Ruby on Rails Tutorial* can be found on the book's website at http://www.railstutorial.org/. If you are reading this book offline, be sure to check the online version of the Rails Tutorial book at http://www.railstutorial.org/book for the latest updates.

sample application will bear more than a passing resemblance to a certain popular social microblogging site (a site which, coincidentally, was also originally written in Rails), the emphasis throughout the tutorial is on general principles, so you will have a solid foundation no matter what kinds of web applications you want to build.

One common question is how much background is necessary to learn web development using the *Ruby on Rails Tutorial*. As discussed in more depth in Section 1.1.1, web development is a challenging subject, especially for complete beginners. Although the tutorial was originally designed for readers with some prior programming or web-development experience, in fact it has found a significant audience among beginning developers. In acknowledgment of this, the present third edition of the *Rails Tutorial* has taken several important steps toward lowering the barrier to getting started with Rails (Box 1.1).

Box 1.1. Lowering the barrier

This third edition of the *Ruby on Rails Tutorial* aims to lower the barrier to getting started with Rails in a number of ways:

- Use of a standard development environment in the cloud (Section 1.2), which sidesteps many of the problems associated with installing and configuring a new system

- Use of the Rails "default stack", including the built-in minitest testing framework

- Elimination of many external dependencies (RSpec, Cucumber, Capybara, Factory Girl)

- A lighter-weight and more flexible approach to testing

- Deferral or elimination of more complex configuration options (Spork, RubyTest)

> - Less emphasis on features specific to any given version of Rails, with greater emphasis on general principles of web development
>
> It is my hope that these changes will make the third edition of the *Ruby on Rails Tutorial* accessible to an even broader audience than previous versions.

In this first chapter, we'll get started with Ruby on Rails by installing all the necessary software and by setting up our development environment (Section 1.2). We'll then create our first Rails application, called `hello_app`. The *Rails Tutorial* emphasizes good software development practices, so immediately after creating our fresh new Rails project we'll put it under version control with Git (Section 1.4). And, believe it or not, in this chapter we'll even put our first app on the wider web by *deploying* it to production (Section 1.5).

In Chapter 2, we'll make a second project, whose purpose is to demonstrate the basic workings of a Rails application. To get up and running quickly, we'll build this *toy app* (called `toy_app`) using scaffolding (Box 1.2) to generate code; because this code is both ugly and complex, Chapter 2 will focus on interacting with the toy app through its *URIs* (often called *URLs*)[2] using a web browser.

The rest of the tutorial focuses on developing a single large *real sample application* (called `sample_app`), writing all the code from scratch. We'll develop the sample app using a combination of *mockups*, *test-driven development* (TDD), and *integration tests*. We'll get started in Chapter 3 by creating static pages and then add a little dynamic content. We'll take a quick detour in Chapter 4 to learn a little about the Ruby language underlying Rails. Then, in Chapter 5 through Chapter 10, we'll complete the foundation for the sample application by making a site layout, a user data model, and a full registration and authentication system (including account activation and password resets). Finally, in Chapter 11 and Chapter 12 we'll add microblogging and social features to make a working example site.

---

[2] *URI* stands for Uniform Resource Identifier, while the slightly less general *URL* stands for Uniform Resource Locator. In practice, the URL is usually equivalent to "the thing you see in the address bar of your browser".

Box 1.2. Scaffolding: Quicker, easier, more seductive

From the beginning, Rails has benefited from a palpable sense of excitement, starting with the famous 15-minute weblog video by Rails creator David Heinemeier Hansson. That video and its successors are a great way to get a taste of Rails' power, and I recommend watching them. But be warned: they accomplish their amazing fifteen-minute feat using a feature called *scaffolding*, which relies heavily on *generated code*, magically created by the Rails **generate scaffold** command.

When writing a Ruby on Rails tutorial, it is tempting to rely on the scaffolding approach—it's quicker, easier, more seductive. But the complexity and sheer amount of code in the scaffolding can be utterly overwhelming to a beginning Rails developer; you may be able to use it, but you probably won't understand it. Following the scaffolding approach risks turning you into a virtuoso script generator with little (and brittle) actual knowledge of Rails.

In the *Ruby on Rails Tutorial*, we'll take the (nearly) polar opposite approach: although Chapter 2 will develop a small toy app using scaffolding, the core of the *Rails Tutorial* is the sample app, which we'll start writing in Chapter 3. At each stage of developing the sample application, we will write *small, bite-sized* pieces of code—simple enough to understand, yet novel enough to be challenging. The cumulative effect will be a deeper, more flexible knowledge of Rails, giving you a good background for writing nearly any type of web application.

1.1   Introduction

Ruby on Rails (or just "Rails" for short) is a web development framework written in the Ruby programming language. Since its debut in 2004, Ruby on Rails has rapidly become one of the most powerful and popular tools for building dynamic web applications. Rails is used by companies as diverse as Airbnb, Basecamp, Disney, GitHub, Hulu, Kickstarter, Shopify, Twitter, and the Yel-

low Pages. There are also many web development shops that specialize in Rails, such as ENTP, thoughtbot, Pivotal Labs, Hashrocket, and HappyFun-Corp, plus innumerable independent consultants, trainers, and contractors.

What makes Rails so great? First of all, Ruby on Rails is 100% open-source, available under the permissive MIT License, and as a result it also costs nothing to download or use. Rails also owes much of its success to its elegant and compact design; by exploiting the malleability of the underlying Ruby language, Rails effectively creates a domain-specific language for writing web applications. As a result, many common web programming tasks—such as generating HTML, making data models, and routing URLs—are easy with Rails, and the resulting application code is concise and readable.

Rails also adapts rapidly to new developments in web technology and framework design. For example, Rails was one of the first frameworks to fully digest and implement the REST architectural style for structuring web applications (which we'll be learning about throughout this tutorial). And when other frameworks develop successful new techniques, Rails creator David Heine-meier Hansson and the Rails core team don't hesitate to incorporate their ideas. Perhaps the most dramatic example is the merger of Rails and Merb, a rival Ruby web framework, so that Rails now benefits from Merb's modular design, stable API, and improved performance.

Finally, Rails benefits from an unusually enthusiastic and diverse community. The results include hundreds of open-source contributors, well-attended conferences, a huge number of gems (self-contained solutions to specific problems such as pagination and image upload), a rich variety of informative blogs, and a cornucopia of discussion forums and IRC channels. The large number of Rails programmers also makes it easier to handle the inevitable application errors: the "Google the error message" algorithm nearly always produces a relevant blog post or discussion-forum thread.

1.1.1 Prerequisites

There are no formal prerequisites to this book—the *Ruby on Rails Tutorial* contains integrated tutorials not only for Rails, but also for the underlying Ruby language, the default Rails testing framework (minitest), the Unix command

line, HTML, CSS, a small amount of JavaScript, and even a little SQL. That's a lot of material to absorb, though, and I generally recommend having some HTML and programming background before starting this tutorial. That said, a surprising number of beginners have used the *Ruby on Rails Tutorial* to learn web development from scratch, so even if you have limited experience I suggest giving it a try. If you feel overwhelmed, you can always go back and start with one of the resources listed below. Another strategy recommended by multiple readers is simply to do the tutorial twice; you may be surprised at how much you learned the first time (and how much easier it is the second time through).

I'm currently working on a series of tutorials representing the full prerequisites to the *Rails Tutorial*, starting with *Learn Enough™ Command Line to Be Dangerous*, an introduction to the Unix command line for complete beginners. You can find a list of subsequent tutorials at learnenough.com, where you can also sign up to be notified when the other tutorials are ready.

One common question when learning Rails is whether to learn Ruby first. The answer depends on your personal learning style and how much programming experience you already have. If you prefer to learn everything systematically from the ground up, or if you have never programmed before, then learning Ruby first might work well for you, and in this case I recommend *Learn to Program* by Chris Pine and *Beginning Ruby* by Peter Cooper. On the other hand, many beginning Rails developers are excited about making *web* applications, and would rather not wait to finish a whole book on Ruby before ever writing a single web page. In this case, I recommend following the short interactive tutorial at Try Ruby[3] to get a general overview before starting with the *Rails Tutorial*. If you still find this tutorial too difficult, you might try starting with *Learn Ruby on Rails* by Daniel Kehoe or One Month Rails, both of which are geared more toward complete beginners than the *Ruby on Rails Tutorial*.

At the end of this tutorial, no matter where you started, you should be ready for the many more intermediate-to-advanced Rails resources out there. Here are some I particularly recommend:

- Code School: Good interactive online programming courses

---

[3]http://tryruby.org/

- The Turing School of Software & Design: a full-time, 27-week training program in Denver, Colorado, with a $500 discount for Rails Tutorial readers using the code RAILSTUTORIAL500

- Bloc: an online bootcamp with a structured curriculum, personalized mentorship, and a focus on learning through concrete projects. Use the coupon code BLOCLOVESHARTL to get $500 off the enrollment fee.

- Launch School: A good online Rails development bootcamp (includes advanced material)

- Firehose Project: A mentor-driven, online coding bootcamp focused on real-world programming skills like test-driven development, algorithms, and building an advanced web application as part of an agile team. Two-week free intro course.

- Thinkful: An online class that pairs you with a professional engineer as you work through a project-based curriculum

- Pragmatic Studio: Online Ruby and Rails courses from Mike and Nicole Clark. Along with *Programming Ruby* author Dave Thomas, Mike taught the first Rails course I took, way back in 2006.

- RailsApps: A large variety of detailed topic-specific Rails projects and tutorials

- Rails Guides: Topical and up-to-date Rails references

1.1.2 Conventions in this book

The conventions in this book are mostly self-explanatory. In this section, I'll mention some that may not be.

Many examples in this book use command-line commands. For simplicity, all command line examples use a Unix-style command line prompt (a dollar sign), as follows:

```
$ echo "hello, world"
hello, world
```

As mentioned in Section 1.2, I recommend that users of all operating systems (especially Windows) use a cloud development environment (Section 1.2.1), which comes with a built-in Unix (Linux) command line. This is particularly useful because Rails comes with many commands that can be run at the command line. For example, in Section 1.3.2 we'll run a local development web server with the **rails server** command:

```
$ rails server
```

As with the command-line prompt, the *Rails Tutorial* uses the Unix convention for directory separators (i.e., a forward slash **/**). For example, the sample application **production.rb** configuration file appears as follows:

```
config/environments/production.rb
```

This file path should be understood as being relative to the application's root directory, which will vary by system; on the cloud IDE (Section 1.2.1), it looks like this:

```
/home/ubuntu/workspace/sample_app/
```

Thus, the full path to **production.rb** is

```
/home/ubuntu/workspace/sample_app/config/environments/production.rb
```

For brevity, I will typically omit the application path and write just **config/-environments/production.rb**.

The *Rails Tutorial* often shows output from various programs (shell commands, version control status, Ruby programs, etc.). Because of the innumerable small differences between different computer systems, the output you see

may not always agree exactly with what is shown in the text, but this is not cause for concern. In addition, some commands may produce errors depending on your system; rather than attempt the Sisyphean task of documenting all such errors in this tutorial, I will delegate to the "Google the error message" algorithm, which among other things is good practice for real-life software development. If you run into any problems while following the tutorial, I suggest consulting the resources listed in the Rails Tutorial help section.[4]

Because the *Rails Tutorial* covers testing of Rails applications, it is often helpful to know if a particular piece of code causes the test suite to fail (indicated by the color red) or pass (indicated by the color green). For convenience, code resulting in a failing test is thus indicated with **RED**, while code resulting in a passing test is indicated with **GREEN**.

Each chapter in the tutorial includes exercises, the completion of which is optional but recommended. In order to keep the main discussion independent of the exercises, the solutions are not generally incorporated into subsequent code listings. In the rare circumstance that an exercise solution is used subsequently, it is explicitly solved in the main text.

Finally, for convenience the *Ruby on Rails Tutorial* adopts two conventions designed to make the many code samples easier to understand. First, some code listings include one or more highlighted lines, as seen below:

```ruby
class User < ActiveRecord::Base
  validates :name,  presence: true
  validates :email, presence: true
end
```

Such highlighted lines typically indicate the most important new code in the given sample, and often (though not always) represent the difference between the present code listing and previous listings. Second, for brevity and simplicity many of the book's code listings include vertical dots, as follows:

---

[4]http://www.railstutorial.org/#help

```
class User < ActiveRecord::Base
  .
  .
  .
  has_secure_password
end
```

These dots represent omitted code and should not be copied literally.

1.2   Up and running

Even for experienced Rails developers, installing Ruby, Rails, and all the associated supporting software can be an exercise in frustration. Compounding the problem is the multiplicity of environments: different operating systems, version numbers, preferences in text editor and integrated development environment (IDE), etc. Users who already have a development environment installed on their local machine are welcome to use their preferred setup, but (as mentioned in Box 1.1) new users are encouraged to sidestep such installation and configuration issues by using a *cloud integrated development environment*. The cloud IDE runs inside an ordinary web browser and hence works the same across different platforms, which is especially useful for operating systems (such as Windows) on which Rails development has historically been difficult. If, despite the challenges involved, you would still prefer to complete the *Ruby on Rails Tutorial* using a local development environment, I recommend following the instructions at InstallRails.com.[5]

1.2.1   Development environment

Considering various idiosyncratic customizations, there are probably as many development environments as there are Rails programmers. To avoid this complexity, the *Ruby on Rails Tutorial* standardizes on the excellent cloud development environment Cloud9. In particular, for this third edition I am pleased to

---

[5]Even then, Windows users should be warned that the Rails installer recommended by InstallRails is often out of date, and is likely to be incompatible with the present tutorial.

partner with Cloud9 to offer a development environment specifically tailored to the needs of this tutorial. The resulting Rails Tutorial Cloud9 workspace comes pre-configured with most of the software needed for professional-grade Rails development, including Ruby, RubyGems, Git. (Indeed, the only big piece of software we'll install separately is Rails itself, and this is intentional (Section 1.2.2).) The cloud IDE also includes the three essential components needed to develop web applications: a text editor, a filesystem navigator, and a command-line terminal (Figure 1.1). Among other features, the cloud IDE text editor supports the "Find in Files" global search that I consider essential to navigating any large Ruby or Rails project.[6] Finally, even if you decide not to use the cloud IDE exclusively in real life (and I certainly recommend learning other tools as well), it provides an excellent introduction to the general capabilities of text editors and other development tools.

Here are the steps for getting started with the cloud development environment:

1. Sign up for a free account at Cloud9[7]

2. Click on "Go to your Dashboard"

3. Select "Create New Workspace"

4. As shown in Figure 1.2, create a workspace called "rails-tutorial" (*not* "rails_tutorial"), set it to "Private to the people I invite", and select the icon for the Rails Tutorial (*not* the icon for Ruby on Rails)

5. Click "Create"

6. After Cloud9 has finished provisioning the workspace, select it and click "Start editing"

Because using two spaces for indentation is a near-universal convention in Ruby, I also recommend changing the editor to use two spaces instead of the default four. As shown in Figure 1.3, you can do this by clicking the gear

---

[6]For example, to find the definition of a function called `foo`, you can do a global search for "def foo".

[7]https://c9.io/web/sign-up/free

Figure 1.1: The anatomy of the cloud IDE.
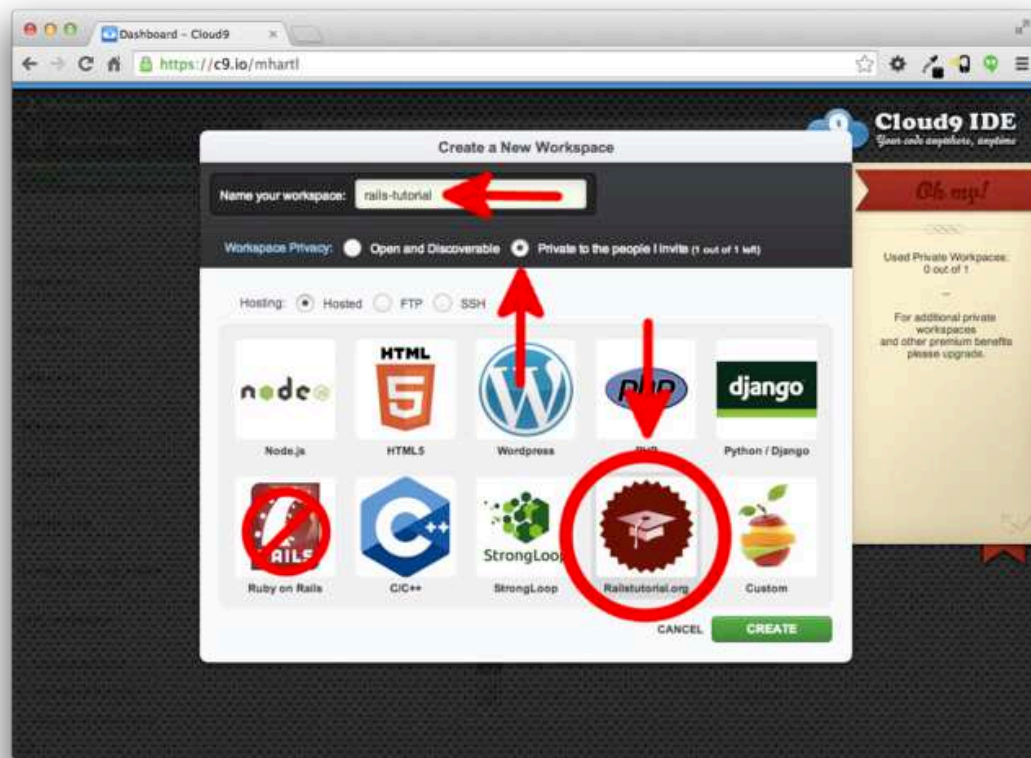
Figure 1.2: Creating a new workspace at Cloud9.

icon in the upper right and then selecting "Code Editor (Ace)" to edit the "Soft Tabs" setting. (Note that this takes effect immediately; you don't need to click a "Save" button.)

1.2.2 Installing Rails

The development environment from Section 1.2.1 includes all the software we need to get started except for Rails itself.[8] To install Rails, we'll use the `gem` command provided by the *RubyGems* package manager, which involves typing

---

[8]At present, Cloud9 includes an older version of Rails that is incompatible with the present tutorial, which is one reason why it's so important to install it ourselves.
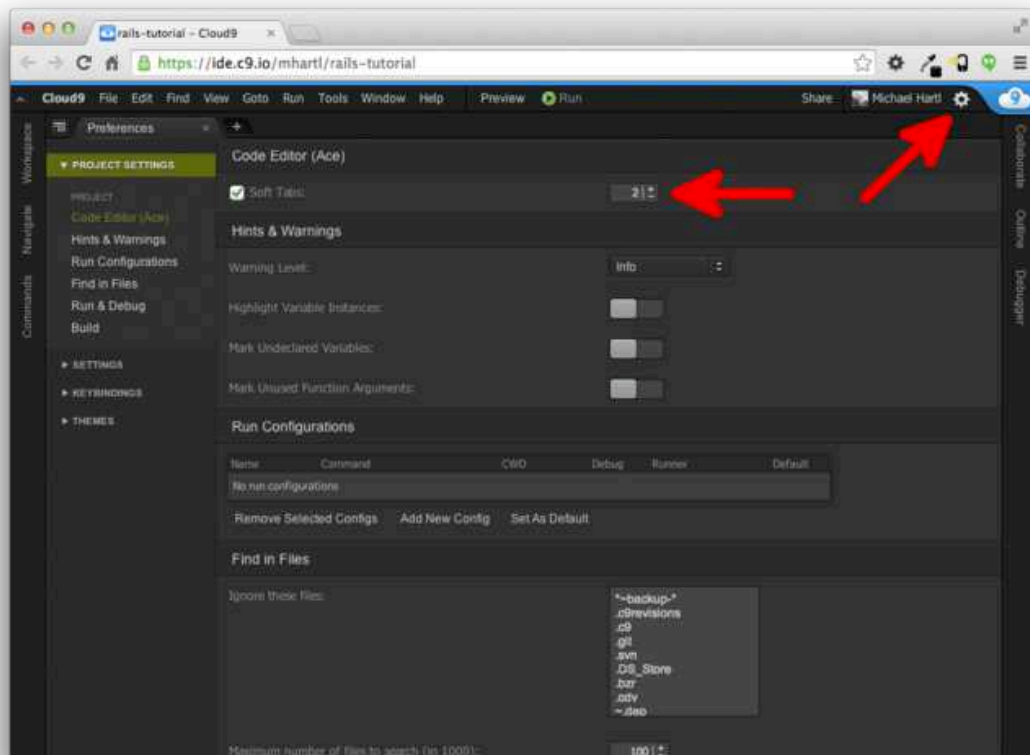
Figure 1.3: Setting Cloud9 to use two spaces for indentation.

the command shown in Listing 1.1 into your command-line terminal. (If developing on your local system, this means using a regular terminal window; if using the cloud IDE, this means using the command-line area shown in Figure 1.1.)

Listing 1.1: Installing Rails with a specific version number.

```
$ gem install rails -v 4.2.2
```

Here the **-v** flag ensures that the specified version of Rails gets installed, which is important to get results consistent with this tutorial.

1.3 The first application

Following a long tradition in computer programming, our goal for the first application is to write a "hello, world" program. In particular, we will create a simple application that displays the string "hello, world!" on a web page, both on our development environment (Section 1.3.4) and on the live web (Section 1.5).

Virtually all Rails applications start the same way, by running the **rails new** command. This handy command creates a skeleton Rails application in a directory of your choice. To get started, users *not* using the Cloud9 IDE recommended in Section 1.2.1 should make a **workspace** directory for your Rails projects if it doesn't already exist (Listing 1.2) and then change into the directory. (Listing 1.2 uses the Unix commands **cd** and **mkdir**; see Box 1.3 if you are not already familiar with these commands.)

Listing 1.2: Making a **workspace** directory for Rails projects (unnecessary in the cloud).

```
$ cd                    # Change to the home directory.
$ mkdir workspace       # Make a workspace directory.
$ cd workspace/         # Change into the workspace directory.
```

Box 1.3. A crash course on the Unix command line

For readers coming from Windows or (to a lesser but still significant extent) Macintosh OS X, the Unix command line may be unfamiliar. Luckily, if you are using the recommended cloud environment, you automatically have access to a Unix (Linux) command line running a standard shell command-line interface known as Bash.

The basic idea of the command line is simple: by issuing short commands, users can perform a large number of operations, such as creating directories (`mkdir`), moving and copying files (`mv` and `cp`), and navigating the filesystem by changing directories (`cd`). Although the command line may seem primitive to users mainly familiar with graphical user interfaces (GUIs), appearances are deceiving: the command line is one of the most powerful tools in the developer's toolbox. Indeed, you will rarely see the desktop of an experienced developer without several open terminal windows running command-line shells.

The general subject is deep, but for the purposes of this tutorial we will need only a few of the most common Unix command-line commands, as summarized in Table 1.1. For a more in-depth treatment of the Unix command line, see *Conquering the Command Line* by Mark Bates (available as a free online version and as ebooks and screencasts).

The next step on both local systems and the cloud IDE is to create the first application using the command in Listing 1.3. Note that Listing 1.3 explicitly includes the Rails version number (`_4.2.2_`) as part of the command. This ensures that the same version of Rails we installed in Listing 1.1 is used to create the first application's file structure. (If the command in Listing 1.3 returns an error like "Could not find 'railties'", it means you don't have the right version of Rails installed, and you should double-check that you followed the command in Listing 1.1 exactly as written.)

Description	Command	Example
list contents	ls	**$ ls -l**
make directory	mkdir <dirname>	**$ mkdir workspace**
change directory	cd <dirname>	**$ cd workspace/**
cd one directory up		**$ cd ..**
cd to home directory		**$ cd ~** or just **$ cd**
cd to path incl. home dir		**$ cd ~/workspace/**
move file (rename)	mv <source> <target>	**$ mv README.rdoc README.md**
copy file	cp <source> <target>	**$ cp README.rdoc README.md**
remove file	rm <file>	**$ rm README.rdoc**
remove empty directory	rmdir <directory>	**$ rmdir workspace/**
remove nonempty directory	rm -rf <directory>	**$ rm -rf tmp/**
concatenate & display file contents	cat <file>	**$ cat ~/.ssh/id_rsa.pub**

Table 1.1: Some common Unix commands.

Listing 1.3: Running **rails new** (with a specific version number).

```
$ cd ~/workspace
$ rails _4.2.2_ new hello_app
      create
      create  README.rdoc
      create  Rakefile
      create  config.ru
      create  .gitignore
      create  Gemfile
      create  app
      create  app/assets/javascripts/application.js
      create  app/assets/stylesheets/application.css
      create  app/controllers/application_controller.rb
      .
      .
      .
      create  test/test_helper.rb
      create  tmp/cache
      create  tmp/cache/assets
      create  vendor/assets/javascripts
      create  vendor/assets/javascripts/.keep
      create  vendor/assets/stylesheets
      create  vendor/assets/stylesheets/.keep
         run  bundle install
Fetching gem metadata from https://rubygems.org/..........
Fetching additional metadata from https://rubygems.org/..
Resolving dependencies...
Using rake 10.3.2
Using i18n 0.6.11
.
```

```
.
.
Your bundle is complete!
Use `bundle show [gemname]` to see where a bundled gem is installed.
       run  bundle exec spring binstub --all
* bin/rake: spring inserted
* bin/rails: spring inserted
```

As seen at the end of Listing 1.3, running **rails new** automatically runs the
bundle install command after the file creation is done. We'll discuss what
this means in more detail starting in Section 1.3.1.

Notice how many files and directories the **rails** command creates. This
standard directory and file structure (Figure 1.4) is one of the many advantages
of Rails; it immediately gets you from zero to a functional (if minimal) ap-
plication. Moreover, since the structure is common to all Rails apps, you can
immediately get your bearings when looking at someone else's code. A sum-
mary of the default Rails files appears in Table 1.2; we'll learn about most of
these files and directories throughout the rest of this book. In particular, start-
ing in Section 5.2.1 we'll discuss the **app/assets** directory, part of the *asset
pipeline* that makes it easier than ever to organize and deploy assets such as
cascading style sheets and JavaScript files.

1.3.1   Bundler

After creating a new Rails application, the next step is to use *Bundler* to in-
stall and include the gems needed by the app. As noted briefly in Section 1.3,
Bundler is run automatically (via **bundle install**) by the **rails** command,
but in this section we'll make some changes to the default application gems
and run Bundler again. This involves opening the **Gemfile** with a text editor.
(With the cloud IDE, this involves clicking the arrow in the file navigator to
open the sample app directory and double-clicking the **Gemfile** icon.) Al-
though the exact version numbers and details may differ slightly, the results
should look something like Figure 1.5 and Listing 1.4. (The code in this file is
Ruby, but don't worry at this point about the syntax; Chapter 4 will cover Ruby
in more depth.) If the files and directories don't appear as shown in Figure 1.5,
click on the file navigator's gear icon and select "Refresh File Tree". (As a
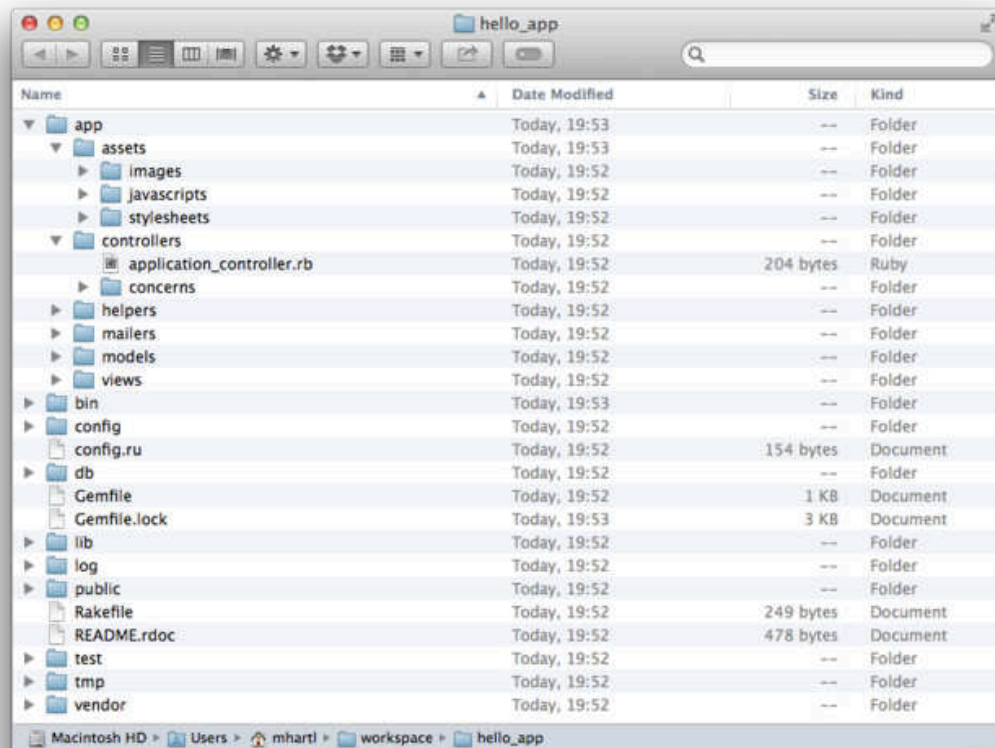
Figure 1.4: The directory structure for a newly created Rails app.

File/Directory	Purpose
`app/`	Core application (app) code, including models, views, controllers, and helpers
`app/assets`	Applications assets such as cascading style sheets (CSS), JavaScript files, and images
`bin/`	Binary executable files
`config/`	Application configuration
`db/`	Database files
`doc/`	Documentation for the application
`lib/`	Library modules
`lib/assets`	Library assets such as cascading style sheets (CSS), JavaScript files, and images
`log/`	Application log files
`public/`	Data accessible to the public (e.g., via web browsers), such as error pages
`bin/rails`	A program for generating code, opening console sessions, or starting a local server
`test/`	Application tests
`tmp/`	Temporary files
`vendor/`	Third-party code such as plugins and gems
`vendor/assets`	Third-party assets such as cascading style sheets (CSS), JavaScript files, and images
`README.rdoc`	A brief description of the application
`Rakefile`	Utility tasks available via the `rake` command
`Gemfile`	Gem requirements for this app
`Gemfile.lock`	A list of gems used to ensure that all copies of the app use the same gem versions
`config.ru`	A configuration file for Rack middleware
`.gitignore`	Patterns for files that should be ignored by Git

Table 1.2: A summary of the default Rails directory structure.
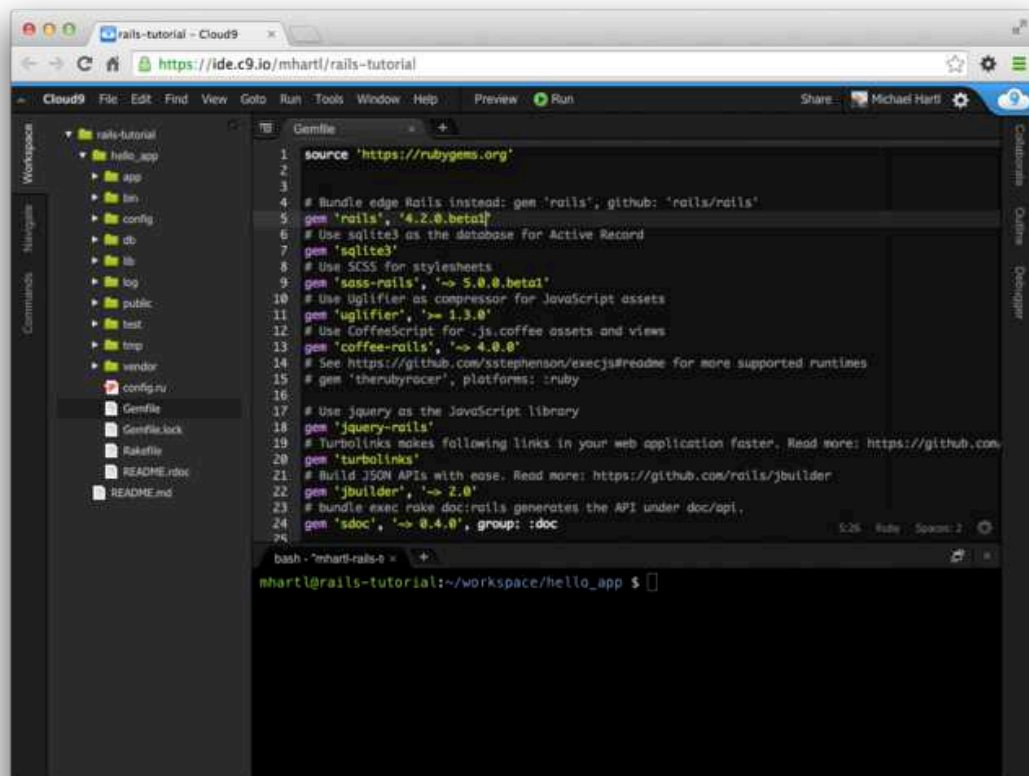
Figure 1.5: The default `Gemfile` open in a text editor.

general rule, you should refresh the file tree any time files or directories don't appear as expected.)

Listing 1.4: The default `Gemfile` in the `hello_app` directory.

```
source 'https://rubygems.org'


# Bundle edge Rails instead: gem 'rails', github: 'rails/rails'
gem 'rails', '4.2.2'
# Use sqlite3 as the database for Active Record
gem 'sqlite3'
# Use SCSS for stylesheets
gem 'sass-rails', '~> 5.0'
```

```ruby
# Use Uglifier as compressor for JavaScript assets
gem 'uglifier', '>= 1.3.0'
# Use CoffeeScript for .js.coffee assets and views
gem 'coffee-rails', '~> 4.1.0'
# See https://github.com/sstephenson/execjs#readme for more supported runtimes
# gem 'therubyracer', platforms: :ruby

# Use jquery as the JavaScript library
gem 'jquery-rails'
# Turbolinks makes following links in your web application faster. Read more:
# https://github.com/rails/turbolinks
gem 'turbolinks'
# Build JSON APIs with ease. Read more: https://github.com/rails/jbuilder
gem 'jbuilder', '~> 2.0'
# bundle exec rake doc:rails generates the API under doc/api.
gem 'sdoc', '~> 0.4.0', group: :doc

# Use ActiveModel has_secure_password
# gem 'bcrypt', '~> 3.1.7'

# Use Unicorn as the app server
# gem 'unicorn'

# Use Capistrano for deployment
# gem 'capistrano-rails', group: :development

group :development, :test do
  # Call 'debugger' anywhere in the code to stop execution and get a
  # debugger console
  gem 'byebug'

  # Access an IRB console on exceptions page and /console in development
  gem 'web-console', '~> 2.0'

  # Spring speeds up development by keeping your application running in the
  # background. Read more: https://github.com/rails/spring
  gem 'spring'
end
```

Many of these lines are commented out with the hash symbol `#`; they are there to show you some commonly needed gems and to give examples of the Bundler syntax. For now, we won't need any gems other than the defaults.

Unless you specify a version number to the `gem` command, Bundler will automatically install the latest requested version of the gem. This is the case, for example, in the code

```
gem 'sqlite3'
```

There are also two common ways to specify a gem version range, which allows us to exert some control over the version used by Rails. The first looks like this:

```
gem 'uglifier', '>= 1.3.0'
```

This installs the latest version of the **uglifier** gem (which handles file compression for the asset pipeline) as long as it's greater than or equal to version **1.3.0**—even if it's, say, version **7.2**. The second method looks like this:

```
gem 'coffee-rails', '~> 4.0.0'
```

This installs the gem **coffee-rails** as long as it's newer than version **4.0.0** and *not* newer than **4.1**. In other words, the >= notation always installs the latest gem, whereas the ~> 4.0.0 notation only installs updated gems representing minor point releases (e.g., from **4.0.0** to **4.0.1**), but not major point releases (e.g., from **4.0** to **4.1**). Unfortunately, experience shows that even minor point releases can break things, so for the *Ruby on Rails Tutorial* we'll err on the side of caution by including exact version numbers for all gems. You are welcome to use the most up-to-date version of any gem, including using the ~> construction in the **Gemfile** (which I generally recommend for more advanced users), but be warned that this may cause the tutorial to act unpredictably.

Converting the **Gemfile** in Listing 1.4 to use exact gem versions results in the code shown in Listing 1.5. (You can determine the exact version number for each gem by running **gem list <gem name>** at the command line, but Listing 1.5 saves you the trouble.) Note that we've also taken this opportunity to arrange for the sqlite3 gem to be included only in a development or test environment (Section 7.1.1), which prevents potential conflicts with the

database used by Heroku (Section 1.5). **Important note: Unless you are reading this at railstutorial.org, you should use the application Gemfile listed at gemfiles-3rd-ed.railstutorial.org instead of the one listed here.**

Listing 1.5: A `Gemfile` with an explicit version for each Ruby gem.

```
source 'https://rubygems.org'

gem 'rails',             '4.2.2'
gem 'sass-rails',        '5.0.2'
gem 'uglifier',          '2.5.3'
gem 'coffee-rails',      '4.1.0'
gem 'jquery-rails',      '4.0.3'
gem 'turbolinks',        '2.3.0'
gem 'jbuilder',          '2.2.3'
gem 'sdoc',              '0.4.0', group: :doc

group :development, :test do
  gem 'sqlite3',     '1.3.9'
  gem 'byebug',      '3.4.0'
  gem 'web-console', '2.0.0.beta3'
  gem 'spring',      '1.1.3'
end
```

Once you've placed the contents of Listing 1.5 into the application's `Gemfile`, install the gems using `bundle install`:[9]

```
$ cd hello_app/
$ bundle install
Fetching source index for https://rubygems.org/
.
.
.
```

The `bundle install` command might take a few moments, but when it's done our application will be ready to run.

---

[9]As noted in Table 3.1, you can even leave off `install`, as the `bundle` command by itself is an alias for `bundle install`.

1.3.2 `rails server`

Thanks to running **`rails new`** in Section 1.3 and **`bundle install`** in Section 1.3.1, we already have an application we can run—but how? Happily, Rails comes with a command-line program, or *script*, that runs a *local* web server to assist us in developing our application. The exact command depends on the environment you're using: on a local system, you just run **`rails server`** (Listing 1.6), whereas on Cloud9 you need to supply an additional *IP binding address* and *port number* to tell the Rails server the address it can use to make the application visible to the outside world (Listing 1.7).[10] (Cloud9 uses the special *environment variables* **`$IP`** and **`$PORT`** to assign the IP address and port number dynamically. If you want to see the values of these variables, type **`echo $IP`** or **`echo $PORT`** at the command line.) If your system complains about the lack of a JavaScript runtime, visit the execjs page at GitHub for a list of possibilities. I particularly recommend installing Node.js.

Listing 1.6: Running the Rails server on a local machine.

```
$ cd ~/workspace/hello_app/
$ rails server
=> Booting WEBrick
=> Rails application starting on http://localhost:3000
=> Run `rails server -h` for more startup options
=> Ctrl-C to shutdown server
```

Listing 1.7: Running the Rails server on the cloud IDE.

```
$ cd ~/workspace/hello_app/
$ rails server -b $IP -p $PORT
=> Booting WEBrick
=> Rails application starting on http://0.0.0.0:8080
=> Run `rails server -h` for more startup options
=> Ctrl-C to shutdown server
```

Whichever option you choose, I recommend running the **`rails server`** command in a second terminal tab so that you can still issue commands in the

---

[10]Normally, websites run on port 80, but this usually requires special privileges, so it's conventional to use a less restricted higher-numbered port for the development server.
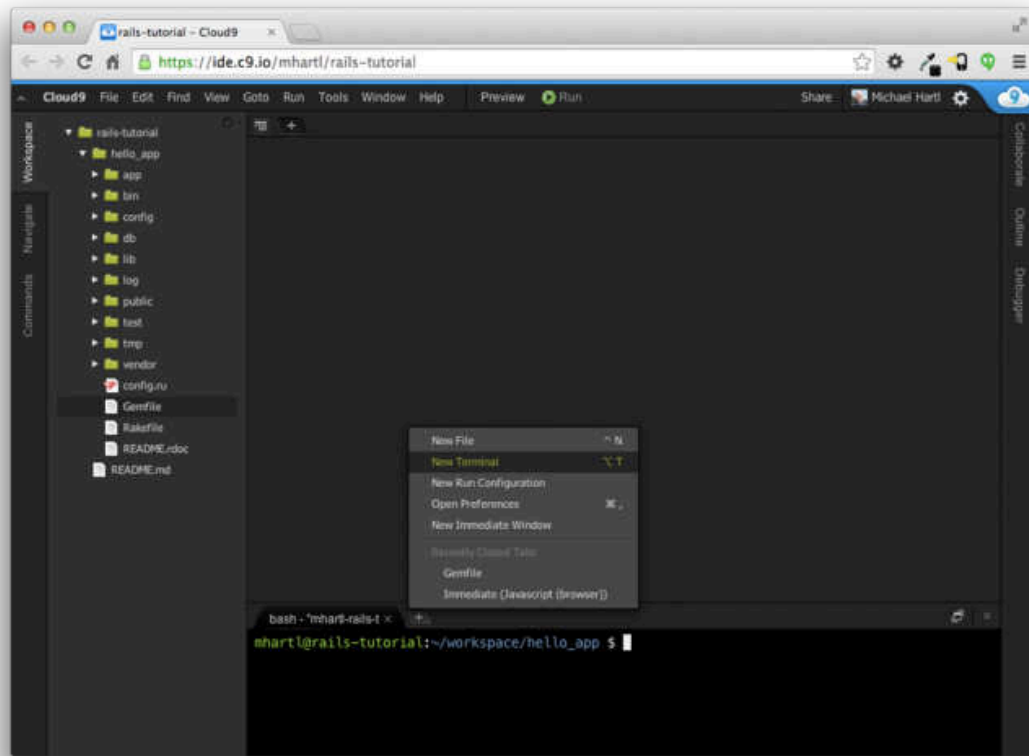
Figure 1.6: Opening a new terminal tab.

first tab, as shown in Figure 1.6 and Figure 1.7. (If you already started a server in your first tab, press Ctrl-C to shut it down.)[11] On a local server, point your browser at the address http://localhost:3000/; on the cloud IDE, go to Share and click on the Application address to open it (Figure 1.8). In either case, the result should look something like Figure 1.9.

   To see information about the first application, click on the link "About your application's environment". Although exact version numbers may differ, the result should look something like Figure 1.10. Of course, we don't need the default Rails page in the long run, but it's nice to see it working for now. We'll

---

[11]Here "C" refers to the character on the keyboard, not the capital letter, so there's no need to hold down the Shift key to get a capital "C".
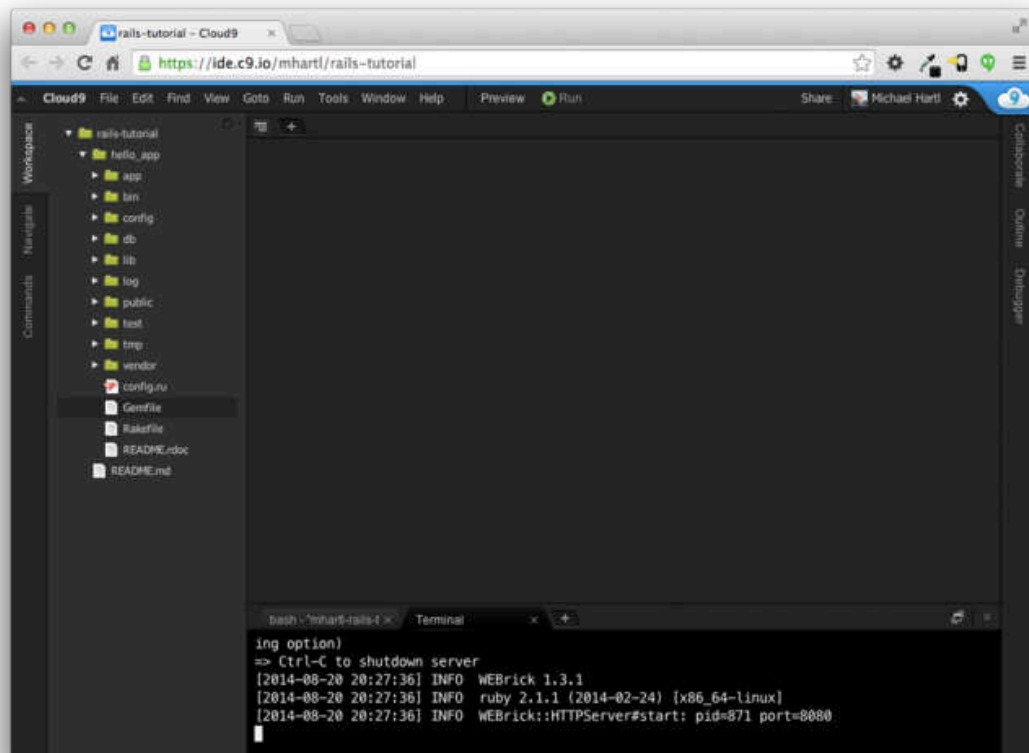
Figure 1.7: Running the Rails server in a separate tab.
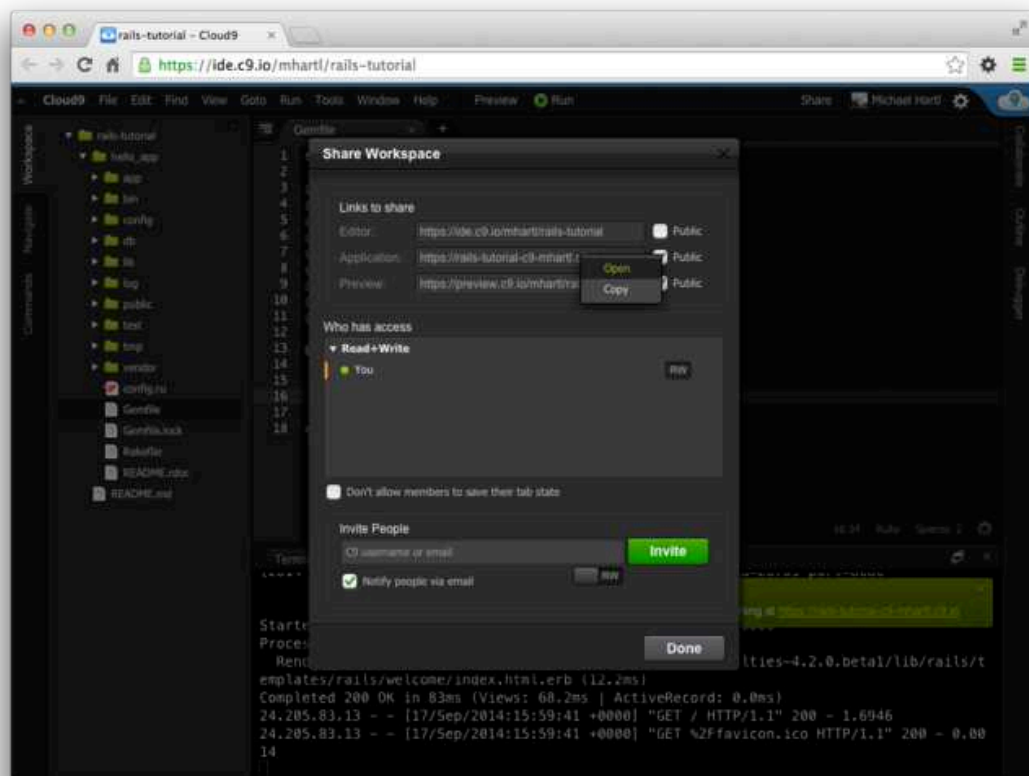
Figure 1.8: Sharing the local server running on the cloud workspace.
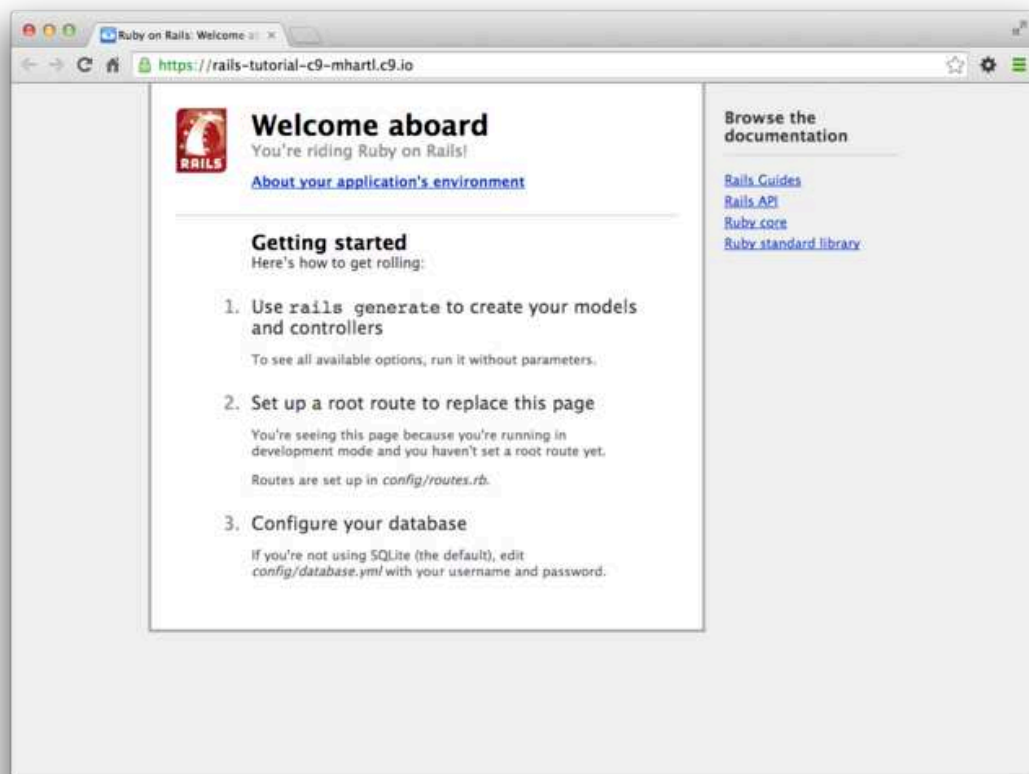
Figure 1.9: The default Rails page served by `rails server`.
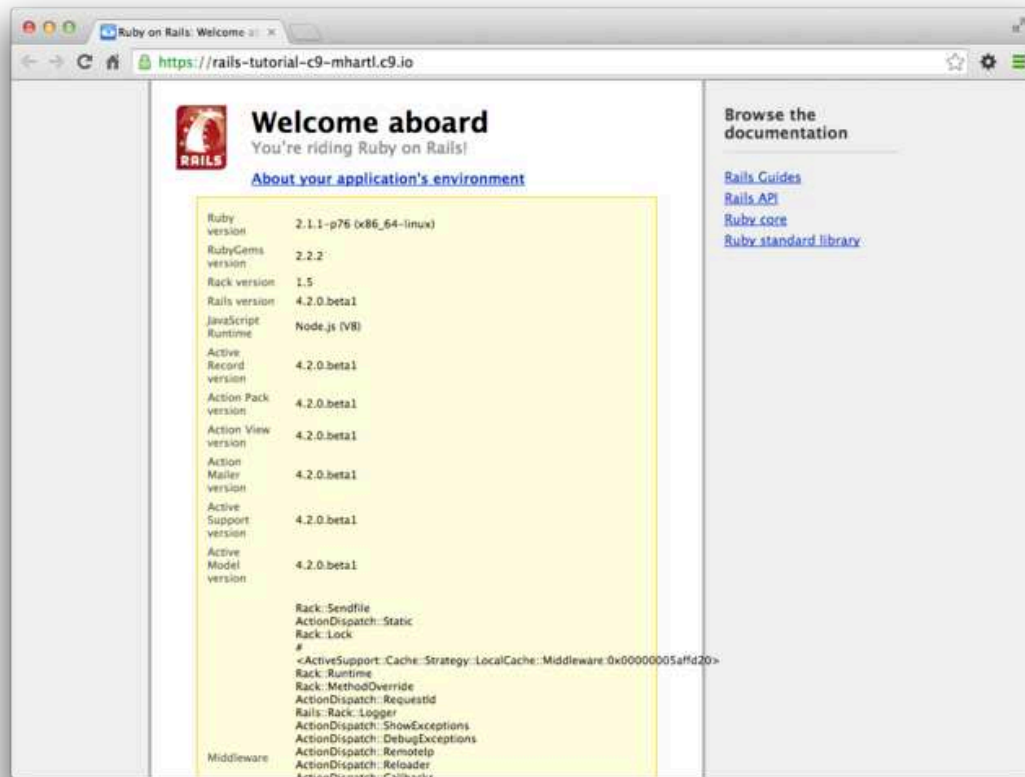
Figure 1.10: The default page with the application's environment.

remove the default page (and replace it with a custom home page) in Section 1.3.4.

1.3.3   Model-View-Controller (MVC)

Even at this early stage, it's helpful to get a high-level overview of how Rails applications work (Figure 1.11). You might have noticed that the standard Rails application structure (Figure 1.4) has an application directory called **app/** with three subdirectories: **models**, **views**, and **controllers**. This is a hint that Rails follows the model-view-controller (MVC) architectural pattern, which

enforces a separation between "domain logic" (also called "business logic") from the input and presentation logic associated with a graphical user interface (GUI). In the case of web applications, the "domain logic" typically consists of data models for things like users, articles, and products, and the GUI is just a web page in a web browser.

When interacting with a Rails application, a browser sends a *request*, which is received by a web server and passed on to a Rails *controller*, which is in charge of what to do next. In some cases, the controller will immediately render a *view*, which is a template that gets converted to HTML and sent back to the browser. More commonly for dynamic sites, the controller interacts with a *model*, which is a Ruby object that represents an element of the site (such as a user) and is in charge of communicating with the database. After invoking the model, the controller then renders the view and returns the complete web page to the browser as HTML.

If this discussion seems a bit abstract right now, worry not; we'll refer back to this section frequently. Section 1.3.4 shows a first tentative application of MVC, while Section 2.2.2 includes a more detailed discussion of MVC in the context of the toy app. Finally, the sample app will use all aspects of MVC; we'll cover controllers and views starting in Section 3.2, models starting in Section 6.1, and we'll see all three working together in Section 7.1.2.

1.3.4 Hello, world!

As a first application of the MVC framework, we'll make a wafer-thin change to the first app by adding a *controller action* to render the string "hello, world!". (We'll learn more about controller actions starting in Section 2.2.2.) The result will be to replace the default Rails page from Figure 1.9 with the "hello, world" page that is the goal of this section.

As implied by their name, controller actions are defined inside controllers. We'll call our action `hello` and place it in the Application controller. Indeed, at this point the Application controller is the only controller we have, which you can verify by running

Figure 1.11: A schematic representation of the model-view-controller (MVC) architecture.

```
$ ls app/controllers/*_controller.rb
```

to view the current controllers. (We'll start creating our own controllers in Chapter 2.) Listing 1.8 shows the resulting definition of **hello**, which uses the **render** function to return the text "hello, world!". (Don't worry about the Ruby syntax right now; it will be covered in more depth in Chapter 4.)

Listing 1.8: Adding a **hello** action to the Application controller.
app/controllers/application_controller.rb

```ruby
class ApplicationController < ActionController::Base
  # Prevent CSRF attacks by raising an exception.
  # For APIs, you may want to use :null_session instead.
  protect_from_forgery with: :exception

  def hello
    render text: "hello, world!"
  end
end
```

Having defined an action that returns the desired string, we need to tell Rails to use that action instead of the default page in Figure 1.10. To do this, we'll edit the Rails *router*, which sits in front of the controller in Figure 1.11 and determines where to send requests that come in from the browser. (I've omitted the router from Figure 1.11 for simplicity, but we'll discuss the router in more detail starting in Section 2.2.2.) In particular, we want to change the default page, the *root route*, which determines the page that is served on the *root URL*. Because it's the URL for an address like http://www.example.com/ (where nothing comes after the final forward slash), the root URL is often referred to as / ("slash") for short.

As seen in Listing 1.9, the Rails routes file (**config/routes.rb**) includes a commented-out line that shows how to structure the root route. Here "welcome" is the controller name and "index" is the action within that controller. To activate the root route, uncomment this line by removing the hash character and then replace it with the code in Listing 1.10, which tells Rails to send the root route to the **hello** action in the Application controller. (As noted

in Section 1.1.2, vertical dots indicate omitted code and should not be copied literally.)

Listing 1.9: The default (commented-out) root route.
config/routes.rb

```
Rails.application.routes.draw do
  .
  .
  .
  # You can have the root of your site routed with "root"
  # root 'welcome#index'
  .
  .
  .
end
```

Listing 1.10: Setting the root route.
config/routes.rb

```
Rails.application.routes.draw do
  .
  .
  .
  # You can have the root of your site routed with "root"
  root 'application#hello'
  .
  .
  .
end
```

With the code from Listing 1.8 and Listing 1.10, the root route returns "hello, world!" as required (Figure 1.12).

1.4   Version control with Git

Now that we have a fresh and working Rails application, we'll take a moment for a step that, while technically optional, would be viewed by experienced software developers as practically essential: placing our application source code under *version control*. Version control systems allow us to track changes to

Figure 1.12: Viewing "hello, world!" in the browser.

our project's code, collaborate more easily, and roll back any inadvertent errors (such as accidentally deleting files). Knowing how to use a version control system is a required skill for every professional-grade software developer.

There are many options for version control, but the Rails community has largely standardized on Git, a distributed version control system originally developed by Linus Torvalds to host the Linux kernel. Git is a large subject, and we'll only be scratching the surface in this book, but there are many good free resources online; I especially recommend Bitbucket Getting Started for a short overview and *Pro Git* by Scott Chacon for a book-length introduction. Putting your source code under version control with Git is *strongly* recommended, not only because it's nearly a universal practice in the Rails world, but also because it will allow you to back up and share your code more easily (Section 1.4.3) and deploy your application right here in the first chapter (Section 1.5).

1.4.1   Installation and setup

The cloud IDE recommended in Section 1.2.1 includes Git by default, so no installation is necessary in this case. Otherwise, InstallRails.com (Section 1.2) includes instructions for installing Git on your system.

First-time system setup

Before using Git, you should perform a set of one-time setup steps. These are *system* setups, meaning you only have to do them once per computer:

```
$ git config --global user.name "Your Name"
$ git config --global user.email your.email@example.com
$ git config --global push.default matching
$ git config --global alias.co checkout
```

Note that the name and email address you use in your Git configuration will be available in any repositories you make public. (Only the first two lines above are strictly necessary. The third line is included only to ensure forward-compatibility with an upcoming release of Git. The optional fourth line is

included so that you can use **co** in place of the more verbose **checkout** command. For maximum compatibility with systems that don't have **co** configured, this tutorial will use the full **checkout** command, but in real life I nearly always use **git co**.)

First-time repository setup

Now we come to some steps that are necessary each time you create a new *repository* (sometimes called a *repo* for short). First navigate to the root directory of the first app and initialize a new repository:

```
$ git init
Initialized empty Git repository in /home/ubuntu/workspace/hello_app/.git/
```

The next step is to add all the project files to the repository using **git add -A**:

```
$ git add -A
```

This command adds all the files in the current directory apart from those that match the patterns in a special file called **.gitignore**. The **rails new** command automatically generates a **.gitignore** file appropriate to a Rails project, but you can add additional patterns as well.[12]

The added files are initially placed in a *staging area*, which contains pending changes to your project. You can see which files are in the staging area using the **status** command:

```
$ git status
On branch master

Initial commit

Changes to be committed:
```

---

[12]Although we'll never need to edit it in the main tutorial, an example of adding a rule to the **.gitignore** file appears in Section 3.7.3, which is part of the optional advanced testing setup in Section 3.7.

```
  (use "git rm --cached <file>..." to unstage)

  new file:    .gitignore
  new file:    Gemfile
  new file:    Gemfile.lock
  new file:    README.rdoc
  new file:    Rakefile
  .
  .
  .
```

(The results are long, so I've used vertical dots to indicate omitted output.)

To tell Git you want to keep the changes, use the **commit** command:

```
$ git commit -m "Initialize repository"
[master (root-commit) df0a62f] Initialize repository
.
.
.
```

The **-m** flag lets you add a message for the commit; if you omit **-m**, Git will open the system's default editor and have you enter the message there. (All the examples in this book will use the **-m** flag.)

It is important to note that Git commits are *local*, recorded only on the machine on which the commits occur. We'll see how to push the changes up to a remote repository (using **git push**) in Section 1.4.4.

By the way, you can see a list of your commit messages using the **log** command:

```
$ git log
commit df0a62f3f091e53ffa799309b3e32c27b0b38eb4
Author: Michael Hartl <michael@michaelhartl.com>
Date:   Wed August 20 19:44:43 2014 +0000

    Initialize repository
```

Depending on the length of your repository's log history, you may have to type **q** to quit.

1.4.2   What good does Git do you?

If you've never used version control before, it may not be entirely clear at this point what good it does you, so let me give just one example. Suppose you've made some accidental changes, such as (D'oh!) deleting the critical **app/controllers/** directory.

```
$ ls app/controllers/
application_controller.rb   concerns/
$ rm -rf app/controllers/
$ ls app/controllers/
ls: app/controllers/: No such file or directory
```

Here we're using the Unix **ls** command to list the contents of the **app/controllers/** directory and the **rm** command to remove it (Table 1.1). The **-rf** flag means "recursive force", which recursively removes all files, directories, subdirectories, and so on, without asking for explicit confirmation of each deletion.

Let's check the status to see what changed:

```
$ git status
On branch master
Changed but not updated:
  (use "git add/rm <file>..." to update what will be committed)
  (use "git checkout -- <file>..." to discard changes in working directory)

      deleted:    app/controllers/application_controller.rb

no changes added to commit (use "git add" and/or "git commit -a")
```

We see here that a file has been deleted, but the changes are only on the "working tree"; they haven't been committed yet. This means we can still undo the changes using the **checkout** command with the **-f** flag to force overwriting the current changes:

```
$ git checkout -f
$ git status
# On branch master
```

```
nothing to commit (working directory clean)
$ ls app/controllers/
application_controller.rb  concerns/
```

The missing files and directories are back. That's a relief!

1.4.3   Bitbucket

Now that we've put our project under version control with Git, it's time to push
our code up to Bitbucket, a site optimized for hosting and sharing Git reposi-
tories. (Previous editions of this tutorial used GitHub instead; see Box 1.4 to
learn the reasons for the switch.) Putting a copy of your Git repository at Bit-
bucket serves two purposes: it's a full backup of your code (including the full
history of commits), and it makes any future collaboration much easier.

Box 1.4. GitHub and Bitbucket

By far the two most popular sites for hosting Git repositories are GitHub and
Bitbucket. The two services share many similarities: both sites allow for Git repos-
itory hosting and collaboration, as well as offering convenient ways to browse and
search repositories. The important differences (from the perspective of this tu-
torial) are that GitHub offers unlimited free repositories (with collaboration) for
open-source repositories while charging for private repos, whereas Bitbucket al-
lows unlimited free private repos while charging for more than a certain number
of collaborators. Which service you use for a particular repo thus depends on your
specific needs.

Previous editions of this book used GitHub because of its emphasis on sup-
porting open-source code, but growing concerns about security have led me to
recommend that *all* web application repositories be private by default. The issue
is that web application repositories might contain potentially sensitive information
such as cryptographic keys and passwords, which could be used to compromise
the security of a site running the code. It is possible, of course, to arrange for this

information to be handled securely (by having Git ignore it, for example), but this is error-prone and requires significant expertise.

As it happens, the sample application created in this tutorial is safe for exposure on the web, but it is dangerous to rely on this fact in general. Thus, to be as secure as possible, we will err on the side of caution and use private repositories by default. Since GitHub charges for private repositories while Bitbucket offers an unlimited number for free, for our purposes Bitbucket is a better fit than GitHub.

Getting started with Bitbucket is simple:

1. Sign up for a Bitbucket account if you don't already have one.

2. Copy your *public key* to your clipboard. As indicated in Listing 1.11, users of the cloud IDE can view their public key using the `cat` command, which can then be selected and copied. If you're using your own system and see no output when running the command in Listing 1.11, follow the instructions on how to install a public key on your Bitbucket account.

3. Add your public key to Bitbucket by clicking on the avatar image in the upper right and selecting "Manage account" and then "SSH keys" (Figure 1.13).

Listing 1.11: Printing the public key using `cat`.

```
$ cat ~/.ssh/id_rsa.pub
```

Once you've added your public key, click on "Create" to create a new repository, as shown in Figure 1.14. When filling in the information for the project, take care to leave the box next to "This is a private repository." checked. After clicking "Create repository", follow the instructions under "Command line > I have an existing project", which should look something like Listing 1.12. (If it doesn't look like Listing 1.12, it might be because the public key didn't get added correctly, in which case I suggest trying that step again.) When pushing
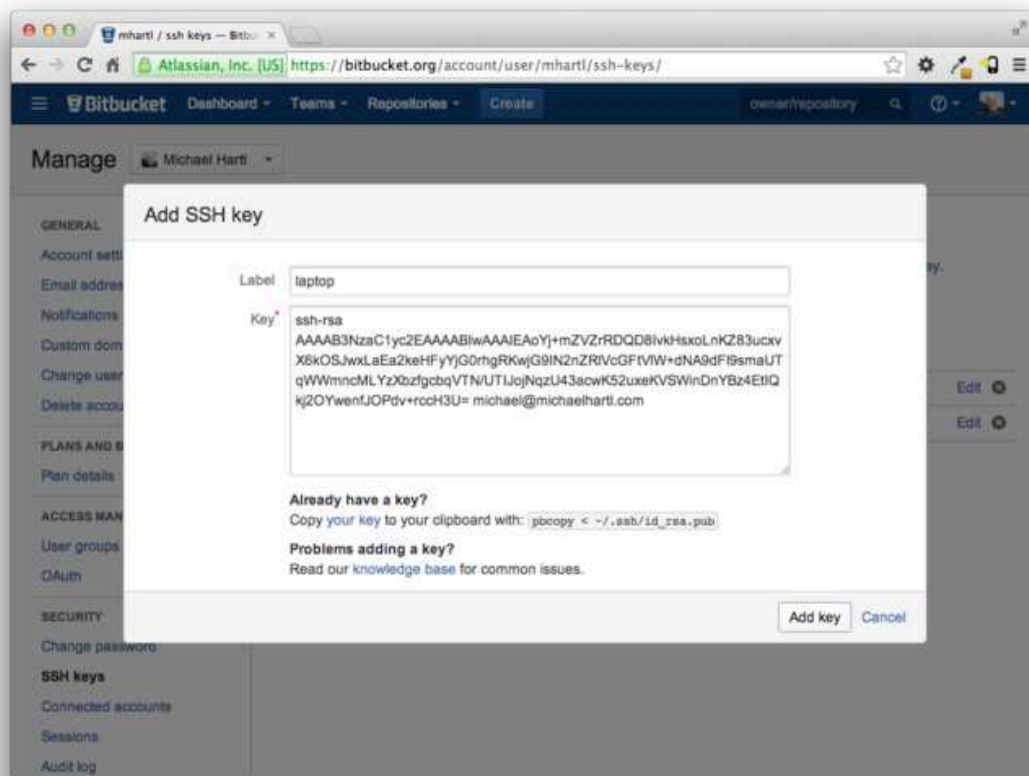
Figure 1.13: Adding the SSH public key.

Figure 1.14: Creating the first app repository at Bitbucket.

up the repository, answer yes if you see the question "Are you sure you want to continue connecting (yes/no)?"

Listing 1.12: Adding Bitbucket and pushing up the repository.

```
$ git remote add origin git@bitbucket.org:<username>/hello_app.git
$ git push -u origin --all # pushes up the repo and its refs for the first time
```

The commands in Listing 1.12 first tell Git that you want to add Bitbucket as the *origin* for your repository, and then push your repository up to the remote origin. (Don't worry about what the **-u** flag does; if you're curious, do a web
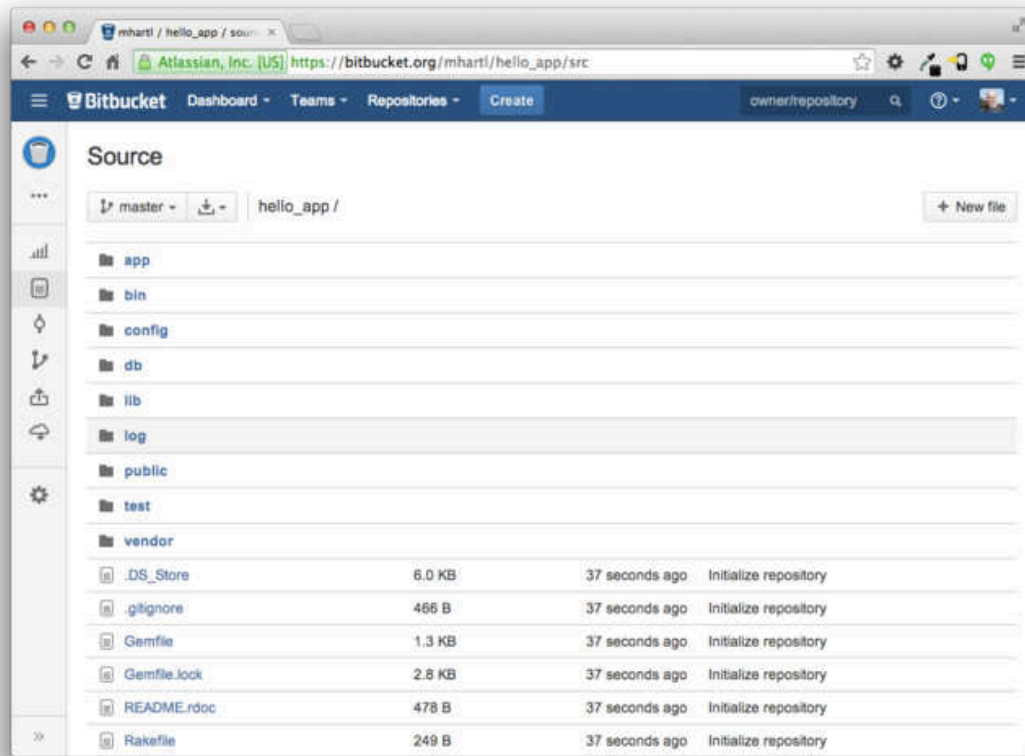
Figure 1.15: A Bitbucket repository page.

search for "git set upstream".) Of course, you should replace `<username>` with your actual username. For example, the command I ran was

```
$ git remote add origin git@bitbucket.org:mhartl/hello_app.git
```

The result is a page at Bitbucket for the hello_app repository, with file browsing, full commit history, and lots of other goodies (Figure 1.15).

1.4.4   Branch, edit, commit, merge

If you've followed the steps in Section 1.4.3, you might notice that Bitbucket didn't automatically detect the **README.rdoc** file from our repository, instead complaining on the main repository page that there is no README present (Figure 1.16). This is an indication that the **rdoc** format isn't common enough for Bitbucket to support it automatically, and indeed I and virtually every other developer I know prefer to use *Markdown* instead. In this section, we'll change the **README.rdoc** file to **README.md**, while taking the opportunity to add some Rails Tutorial–specific content to the README file. In the process, we'll see a first example of the branch, edit, commit, merge workflow that I recommend using with Git.[13]

Branch

Git is incredibly good at making *branches*, which are effectively copies of a repository where we can make (possibly experimental) changes without modifying the parent files. In most cases, the parent repository is the *master* branch, and we can create a new topic branch by using **checkout** with the **-b** flag:

```
$ git checkout -b modify-README
Switched to a new branch 'modify-README'
$ git branch
  master
* modify-README
```

Here the second command, **git branch**, just lists all the local branches, and the asterisk ***** identifies which branch we're currently on. Note that **git checkout -b modify-README** both creates a new branch and switches to it, as indicated by the asterisk in front of the **modify-README** branch. (If you set up the **co** alias in Section 1.4, you can use **git co -b modify-README** instead.)

   The full value of branching only becomes clear when working on a project with multiple developers,[14] but branches are helpful even for a single-developer

---

[13]For a convenient way to visualize Git repositories, take a look at Atlassian's SourceTree app.

[14]See the chapter Git Branching in *Pro Git* for details.
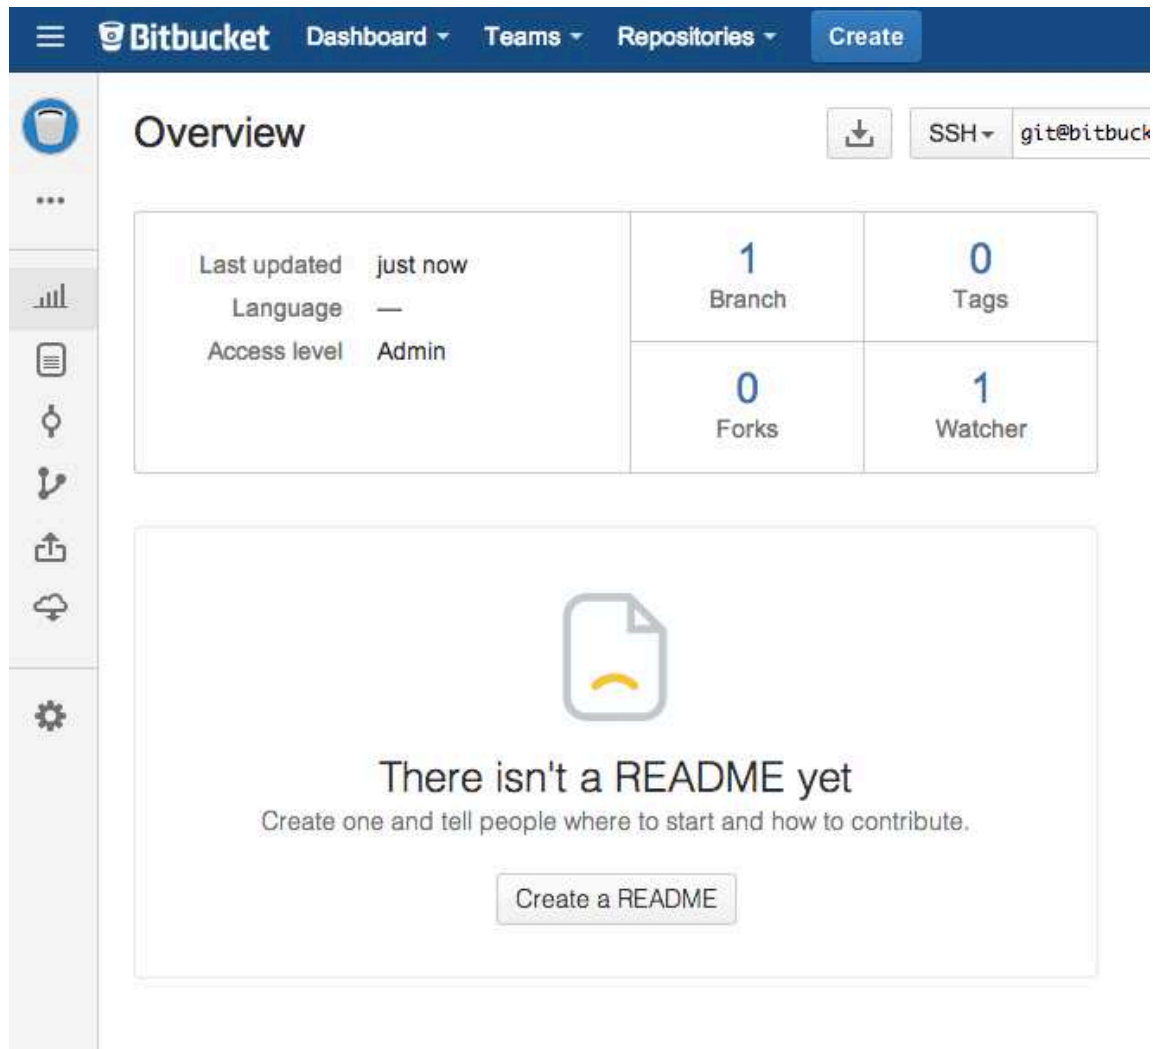
Figure 1.16: Bitbucket's message for a missing README.

tutorial such as this one. In particular, the master branch is insulated from any changes we make to the topic branch, so even if we *really* screw things up we can always abandon the changes by checking out the master branch and deleting the topic branch. We'll see how to do this at the end of the section.

By the way, for a change as small as this one I wouldn't normally bother with a new branch, but in the present context it's a prime opportunity to start practicing good habits.

Edit

After creating the topic branch, we'll edit it to make it a little more descriptive. I prefer the Markdown markup language to the default RDoc for this purpose, and if you use the file extension `.md` then Bitbucket will automatically format it nicely for you. So, first we'll use Git's version of the Unix `mv` (move) command to change the name:

```
$ git mv README.rdoc README.md
```

Then fill `README.md` with the contents of Listing 1.13.

Listing 1.13: The new `README` file, `README.md`.

```
# Ruby on Rails Tutorial: "hello, world!"

This is the first application for the
[*Ruby on Rails Tutorial*](http://www.railstutorial.org/)
by [Michael Hartl](http://www.michaelhartl.com/).
```

Commit

With the changes made, we can take a look at the status of our branch:

```
$ git status
On branch modify-README
Changes to be committed:
```

```
  (use "git reset HEAD <file>..." to unstage)

  renamed:    README.rdoc -> README.md

Changes not staged for commit:
  (use "git add <file>..." to update what will be committed)
  (use "git checkout -- <file>..." to discard changes in working directory)

  modified:   README.md
```

At this point, we could use **git add -A** as in Section 1.4.1, but **git commit** provides the **-a** flag as a shortcut for the (very common) case of committing all modifications to existing files (or files created using **git mv**, which don't count as new files to Git):

```
$ git commit -a -m "Improve the README file"
2 files changed, 5 insertions(+), 243 deletions(-)
delete mode 100644 README.rdoc
create mode 100644 README.md
```

Be careful about using the **-a** flag improperly; if you have added any new files to the project since the last commit, you still have to tell Git about them using **git add -A** first.

Note that we write the commit message in the *present* tense (and, technically speaking, the imperative mood). Git models commits as a series of patches, and in this context it makes sense to describe what each commit *does*, rather than what it did. Moreover, this usage matches up with the commit messages generated by Git commands themselves. See the article "Shiny new commit styles" for more information.

Merge

Now that we've finished making our changes, we're ready to *merge* the results back into our master branch:

```
$ git checkout master
Switched to branch 'master'
$ git merge modify-README
Updating 34f06b7..2c92bef
Fast forward
README.rdoc     |  243 ---------------------------------------------
README.md       |    5 +
2 files changed, 5 insertions(+), 243 deletions(-)
delete mode 100644 README.rdoc
create mode 100644 README.md
```

Note that the Git output frequently includes things like `34f06b7`, which are related to Git's internal representation of repositories. Your exact results will differ in these details, but otherwise should essentially match the output shown above.

After you've merged in the changes, you can tidy up your branches by deleting the topic branch using `git branch -d` if you're done with it:

```
$ git branch -d modify-README
Deleted branch modify-README (was 2c92bef).
```

This step is optional, and in fact it's quite common to leave the topic branch intact. This way you can switch back and forth between the topic and master branches, merging in changes every time you reach a natural stopping point.

As mentioned above, it's also possible to abandon your topic branch changes, in this case with `git branch -D`:

```
# For illustration only; don't do this unless you mess up a branch
$ git checkout -b topic-branch
$ <really screw up the branch>
$ git add -A
$ git commit -a -m "Major screw up"
$ git checkout master
$ git branch -D topic-branch
```

Unlike the `-d` flag, the `-D` flag will delete the branch even though we haven't merged in the changes.

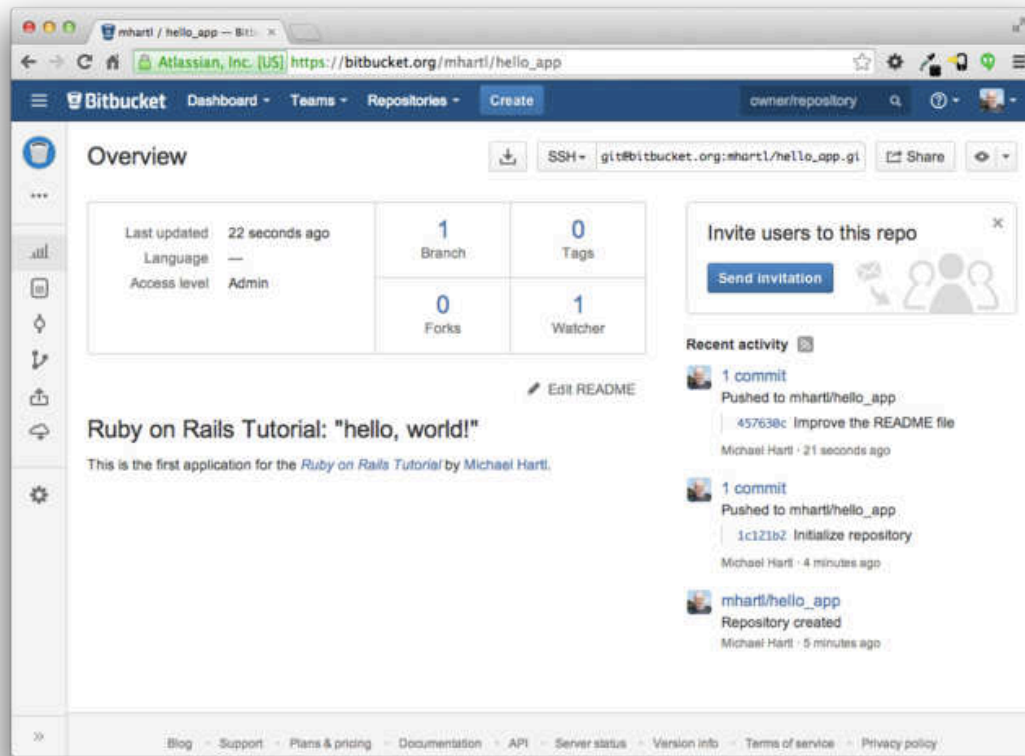Figure 1.17: The improved **README** file formatted with Markdown.

Push

Now that we've updated the **README**, we can push the changes up to Bitbucket
to see the result. Since we have already done one push (Section 1.4.3), on most
systems we can omit **origin master**, and simply run **git push**:

```
$ git push
```

As promised in Section 1.4.4, Bitbucket nicely formats the new file using Mark-
down (Figure 1.17).

1.5 Deploying

Even at this early stage, we're already going to deploy our (nearly empty) Rails application to production. This step is optional, but deploying early and often allows us to catch any deployment problems early in our development cycle. The alternative—deploying only after laborious effort sealed away in a development environment—often leads to terrible integration headaches when launch time comes.[15]

Deploying Rails applications used to be a pain, but the Rails deployment ecosystem has matured rapidly in the past few years, and now there are several great options. These include shared hosts or virtual private servers running Phusion Passenger (a module for the Apache and Nginx[16] web servers), full-service deployment companies such as Engine Yard and Rails Machine, and cloud deployment services such as Engine Yard Cloud, Ninefold, and Heroku.

My favorite Rails deployment option is Heroku, which is a hosted platform built specifically for deploying Rails and other web applications. Heroku makes deploying Rails applications ridiculously easy—as long as your source code is under version control with Git. (This is yet another reason to follow the Git setup steps in Section 1.4 if you haven't already.) In addition, for many purposes, including for this tutorial, Heroku's free tier is more than sufficient. Indeed, the first two editions of this tutorial were hosted for free on Heroku, which served several million requests without charging me a cent.

The rest of this section is dedicated to deploying our first application to Heroku. Some of the ideas are fairly advanced, so don't worry about understanding all the details; what's important is that by the end of the process we'll have deployed our application to the live web.

1.5.1 Heroku setup

Heroku uses the PostgreSQL database (pronounced "post-gres-cue-ell", and often called "Postgres" for short), which means that we need to add the `pg` gem

---

[15]Though it shouldn't matter for the example applications in the *Rails Tutorial*, if you're worried about accidentally making your app public too soon there are several options; see Section 1.5.4 for one.

[16]Pronounced "Engine X".

in the production environment to allow Rails to talk to Postgres:[17]

```
group :production do
  gem 'pg',            '0.17.1'
  gem 'rails_12factor', '0.0.2'
end
```

Note also the addition of the rails_12factor gem, which is used by He-
roku to serve static assets such as images and stylesheets. Finally, be sure
to incorporate the changes made in Listing 1.5 preventing the sqlite3 gem
from being included in a production environment, since SQLite isn't supported
at Heroku:

```
group :development, :test do
  gem 'sqlite3',     '1.3.9'
  gem 'byebug',      '3.4.0'
  gem 'web-console', '2.0.0.beta3'
  gem 'spring',      '1.1.3'
end
```

The resulting **Gemfile** appears as in Listing 1.14.

Listing 1.14: A **Gemfile** with added gems.

```
source 'https://rubygems.org'

gem 'rails',        '4.2.2'
gem 'sass-rails',   '5.0.2'
gem 'uglifier',     '2.5.3'
gem 'coffee-rails', '4.1.0'
gem 'jquery-rails', '4.0.3'
gem 'turbolinks',   '2.3.0'
gem 'jbuilder',     '2.2.3'
gem 'sdoc',         '0.4.0', group: :doc

group :development, :test do
  gem 'sqlite3',     '1.3.9'
  gem 'byebug',      '3.4.0'
```

---

[17]Generally speaking, it's a good idea for the development and production environments to match as closely as
possible, which includes using the same database, but for the purposes of this tutorial we'll always use SQLite
locally and PostgreSQL in production. See Section 3.1 for more information.

```
  gem 'web-console', '2.0.0.beta3'
  gem 'spring',      '1.1.3'
end

group :production do
  gem 'pg',              '0.17.1'
  gem 'rails_12factor', '0.0.2'
end
```

To prepare the system for deployment to production, we run **bundle install** with a special flag to prevent the local installation of any production gems (which in this case consists of pg and rails_12factor):

```
$ bundle install --without production
```

Because the only gems added in Listing 1.14 are restricted to a production environment, right now this command doesn't actually install any additional local gems, but it's needed to update **Gemfile.lock** with the pg and rails_-12factor gems. We can commit the resulting change as follows:

```
$ git commit -a -m "Update Gemfile.lock for Heroku"
```

Next we have to create and configure a new Heroku account. The first step is to sign up for Heroku. Then check to see if your system already has the Heroku command-line client installed:

```
$ heroku version
```

Those using the cloud IDE should see the Heroku version number, indicating that the **heroku** CLI is available, but on other systems it may be necessary to install it using the Heroku Toolbelt.[18]

Once you've verified that the Heroku command-line interface is installed, use the **heroku** command to log in and add your SSH key:

---

[18]https://toolbelt.heroku.com/

```
$ heroku login
$ heroku keys:add
```

Finally, use the **heroku create** command to create a place on the Heroku servers for the sample app to live (Listing 1.15).

Listing 1.15: Creating a new application at Heroku.
```
$ heroku create
Creating damp-fortress-5769... done, stack is cedar
http://damp-fortress-5769.herokuapp.com/ | git@heroku.com:damp-fortress-5769.git
Git remote heroku added
```

The **heroku** command creates a new subdomain just for our application, available for immediate viewing. There's nothing there yet, though, so let's get busy deploying.

1.5.2   Heroku deployment, step one

To deploy the application, the first step is to use Git to push the master branch up to Heroku:

```
$ git push heroku master
```

(You may see some warning messages, which you should ignore for now. We'll discuss them further in Section 7.5.)

1.5.3   Heroku deployment, step two

There is no step two! We're already done. To see your newly deployed application, visit the address that you saw when you ran **heroku create** (i.e., Listing 1.15). (If you're working on your local machine instead of the cloud IDE, you can also use **heroku open**.) The result appears in Figure 1.18. The page is identical to Figure 1.12, but now it's running in a production environment on the live web.

Figure 1.18: The first Rails Tutorial application running on Heroku.

1.5.4   Heroku commands

There are many Heroku commands, and we'll barely scratch the surface in this book. Let's take a minute to show just one of them by renaming the application as follows:

```
$ heroku rename rails-tutorial-hello
```

Don't use this name yourself; it's already taken by me! In fact, you probably shouldn't bother with this step right now; using the default address supplied by Heroku is fine. But if you do want to rename your application, you can arrange for it to be reasonably secure by using a random or obscure subdomain, such as the following:

```
hwpcbmze.herokuapp.com
seyjhflo.herokuapp.com
jhyicevg.herokuapp.com
```

With a random subdomain like this, someone could visit your site only if you gave them the address. (By the way, as a preview of Ruby's compact awesomeness, here's the code I used to generate the random subdomains:

```
('a'..'z').to_a.shuffle[0..7].join
```

Pretty sweet.)

In addition to supporting subdomains, Heroku also supports custom domains. (In fact, the Ruby on Rails Tutorial site lives at Heroku; if you're reading this book online, you're looking at a Heroku-hosted site right now!) See the Heroku documentation for more information about custom domains and other Heroku topics.

1.6 Conclusion

We've come a long way in this chapter: installation, development environment setup, version control, and deployment. In the next chapter, we'll build on the foundation from Chapter 1 to make a database-backed *toy app*, which will give us our first real taste of what Rails can do.

If you'd like to share your progress at this point, feel free to send a tweet or Facebook status update with something like this:

<div align="center">

I'm learning Ruby on Rails with the @railstutorial!
http://www.railstutorial.org/

</div>

I also recommend signing up for the Rails Tutorial email list[19], which will ensure that you receive priority updates (and exclusive coupon codes) regarding the *Ruby on Rails Tutorial*.

1.6.1 What we learned in this chapter

- Ruby on Rails is a web development framework written in the Ruby programming language.

- Installing Rails, generating an application, and editing the resulting files is easy using a pre-configured cloud environment.

- Rails comes with a command-line command called `rails` that can generate new applications (`rails new`) and run local servers (`rails server`).

- We added a controller action and modified the root route to create a "hello, world" application.

- We protected against data loss while enabling collaboration by placing our application source code under version control with Git and pushing the resulting code to a private repository at Bitbucket.

- We deployed our application to a production environment using Heroku.

---

[19]http://www.railstutorial.org/#email

1.7 Exercises

- You can get the solutions manual for the exercises with any purchase at www.railstutorial.org.

- Sign up for the Rails Tutorial email list and get Rails Tutorial cheatsheet cards for the first and toy apps as a free bonus.

1. Change the content of the **hello** action in Listing 1.8 to read "hola, mundo!" instead of "hello, world!". *Extra credit*: Show that Rails supports non-ASCII characters by including an inverted exclamation point, as in "¡Hola, mundo!" (Figure 1.19).[20]

2. By following the example of the **hello** action in Listing 1.8, add a second action called **goodbye** that renders the text "goodbye, world!". Edit the routes file from Listing 1.10 so that the root route goes to **goodbye** instead of to **hello** (Figure 1.20).

---

[20]Your editor may display a message like "invalid multibyte character", but this is not a cause for concern. You can Google the error message if you want to learn how to make it go away.

Figure 1.19: Changing the root route to return "¡Hola, mundo!".

Figure 1.20: Changing the root route to return "goodbye, world!".

Chapter 2

A toy app

In this chapter, we'll develop a toy demo application to show off some of the power of Rails. The purpose is to get a high-level overview of Ruby on Rails programming (and web development in general) by rapidly generating an application using *scaffold generators*, which create a large amount of functionality automatically. As discussed in Box 1.2, the rest of the book will take the opposite approach, developing a full sample application incrementally and explaining each new concept as it arises, but for a quick overview (and some instant gratification) there is no substitute for scaffolding. The resulting toy app will allow us to interact with it through its URLs, giving us insight into the structure of a Rails application, including a first example of the *REST architecture* favored by Rails.

As with the forthcoming sample application, the toy app will consist of *users* and their associated *microposts* (thus constituting a minimalist Twitter-style app). The functionality will be utterly under-developed, and many of the steps will seem like magic, but worry not: the full sample app will develop a similar application from the ground up starting in Chapter 3, and I will provide plentiful forward-references to later material. In the mean time, have patience and a little faith—the whole point of this tutorial is to take you *beyond* this superficial, scaffold-driven approach to achieve a deeper understanding of Rails.

2.1   Planning the application

In this section, we'll outline our plans for the toy application. As in Section 1.3, we'll start by generating the application skeleton using the `rails new` command with a specific Rails version number:

```
$ cd ~/workspace
$ rails _4.2.2_ new toy_app
$ cd toy_app/
```

If the command above returns an error like "Could not find 'railties'", it means you dont have the right version of Rails installed, and you should double-check that you followed the command in Listing 1.1 exactly as written. (If you're using the cloud IDE as recommended in Section 1.2.1, note that this second app can be created in the same workspace as the first. It is not necessary to create a new workspace. In order to get the files to appear, you may need to click the gear icon in the file navigator area and select "Refresh File Tree".)

Next, we'll use a text editor to update the `Gemfile` needed by Bundler with the contents of Listing 2.1. **Important note: Unless you are reading this at railstutorial.org, you should use the application Gemfile listed at gemfiles-3rd-ed.railstutorial.org instead of the one listed here.**

Listing 2.1: A `Gemfile` for the toy app.

```ruby
source 'https://rubygems.org'

gem 'rails',        '4.2.2'
gem 'sass-rails',    '5.0.2'
gem 'uglifier',      '2.5.3'
gem 'coffee-rails', '4.1.0'
gem 'jquery-rails', '4.0.3'
gem 'turbolinks',    '2.3.0'
gem 'jbuilder',      '2.2.3'
gem 'sdoc',          '0.4.0', group: :doc

group :development, :test do
  gem 'sqlite3',      '1.3.9'
  gem 'byebug',       '3.4.0'
  gem 'web-console', '2.0.0.beta3'
  gem 'spring',       '1.1.3'
```

```
end

group :production do
  gem 'pg',           '0.17.1'
  gem 'rails_12factor', '0.0.2'
end
```

Note that Listing 2.1 is identical to Listing 1.14.

As in Section 1.5.1, we'll install the local gems while suppressing the installation of production gems using the `--without production` option:

```
$ bundle install --without production
```

Finally, we'll put the toy app under version control with Git:

```
$ git init
$ git add -A
$ git commit -m "Initialize repository"
```

You should also create a new repository by clicking on the "Create" button at Bitbucket (Figure 2.1), and then push up to the remote repository:

```
$ git remote add origin git@bitbucket.org:<username>/toy_app.git
$ git push -u origin --all # pushes up the repo and its refs for the first time
```

Finally, it's never too early to deploy, which I suggest doing by following the same "hello, world!" steps in Listing 1.8 and Listing 1.9.[1] Then commit the changes and push up to Heroku:

```
$ git commit -am "Add hello"
$ heroku create
$ git push heroku master
```

---

[1] The main reason for this is that the default Rails page typically breaks at Heroku, which makes it hard to tell if the deployment was successful or not.

Figure 2.1: Creating the toy app repository at Bitbucket.

users	
`id`	integer
`name`	string
`email`	string

Figure 2.2: The data model for users.

(As in Section 1.5, you may see some warning messages, which you should ignore for now. We'll eliminate them in Section 7.5.) Apart from the address of the Heroku app, the result should be the same as in Figure 1.18.

Now we're ready to start making the app itself. The typical first step when making a web application is to create a *data model*, which is a representation of the structures needed by our application. In our case, the toy app will be a microblog, with only users and short (micro)posts. Thus, we'll begin with a model for *users* of the app (Section 2.1.1), and then we'll add a model for *microposts* (Section 2.1.2).

2.1.1 A toy model for users

There are as many choices for a user data model as there are different registration forms on the web; we'll go with a distinctly minimalist approach. Users of our toy app will have a unique `integer` identifier called `id`, a publicly viewable `name` (of type `string`), and an `email` address (also a `string`) that will double as a username. A summary of the data model for users appears in Figure 2.2.

As we'll see starting in Section 6.1.1, the label `users` in Figure 2.2 corresponds to a *table* in a database, and the `id`, `name`, and `email` attributes are *columns* in that table.

microposts	
id	integer
content	text
user_id	integer

Figure 2.3: The data model for microposts.

2.1.2   A toy model for microposts

The core of the micropost data model is even simpler than the one for users: a micropost has only an **id** and a **content** field for the micropost's text (of type **text**).[2] There's an additional complication, though: we want to *associate* each micropost with a particular user. We'll accomplish this by recording the **user_id** of the owner of the post. The results are shown in Figure 2.3.

We'll see in Section 2.3.3 (and more fully in Chapter 11) how this **user_id** attribute allows us to succinctly express the notion that a user potentially has many associated microposts.

2.2   The Users resource

In this section, we'll implement the users data model in Section 2.1.1, along with a web interface to that model. The combination will constitute a *Users resource*, which will allow us to think of users as objects that can be created, read, updated, and deleted through the web via the HTTP protocol. As promised in the introduction, our Users resource will be created by a scaffold generator program, which comes standard with each Rails project. I urge you not to look too closely at the generated code; at this stage, it will only serve to confuse you.

---

[2]Because microposts are short by design, the **string** type is actually big enough to contain them, but using **text** better expresses our intent, while also giving us greater flexibility should we ever wish to relax the length constraint.

Rails scaffolding is generated by passing the **scaffold** command to the **rails generate** script. The argument of the **scaffold** command is the singular version of the resource name (in this case, **User**), together with optional parameters for the data model's attributes:[3]

```
$ rails generate scaffold User name:string email:string
      invoke  active_record
      create    db/migrate/20140821011110_create_users.rb
      create    app/models/user.rb
      invoke    test_unit
      create      test/models/user_test.rb
      create      test/fixtures/users.yml
      invoke  resource_route
       route    resources :users
      invoke  scaffold_controller
      create    app/controllers/users_controller.rb
      invoke    erb
      create      app/views/users
      create      app/views/users/index.html.erb
      create      app/views/users/edit.html.erb
      create      app/views/users/show.html.erb
      create      app/views/users/new.html.erb
      create      app/views/users/_form.html.erb
      invoke    test_unit
      create      test/controllers/users_controller_test.rb
      invoke    helper
      create      app/helpers/users_helper.rb
      invoke      test_unit
      create        test/helpers/users_helper_test.rb
      invoke    jbuilder
      create      app/views/users/index.json.jbuilder
      create      app/views/users/show.json.jbuilder
      invoke  assets
      invoke    coffee
      create      app/assets/javascripts/users.js.coffee
      invoke    scss
      create      app/assets/stylesheets/users.css.scss
      invoke  scss
      create    app/assets/stylesheets/scaffolds.css.scss
```

By including **name:string** and **email:string**, we have arranged for the User model to have the form shown in Figure 2.2. (Note that there is no need to include a parameter for **id**; it is created automatically by Rails for use as the *primary key* in the database.)

---

[3]The name of the scaffold follows the convention of *models*, which are singular, rather than resources and controllers, which are plural. Thus, we have **User** instead of **Users**.

To proceed with the toy application, we first need to *migrate* the database using *Rake* (Box 2.1):

```
$ bundle exec rake db:migrate
==  CreateUsers: migrating ======================================================
-- create_table(:users)
   -> 0.0017s
==  CreateUsers: migrated (0.0018s) =========================================
```

This simply updates the database with our new **users** data model. (We'll learn more about database migrations starting in Section 6.1.1.) Note that, in order to ensure that the command uses the version of Rake corresponding to our **Gemfile**, we need to run **rake** using **bundle exec**. On many systems, including the cloud IDE, you can omit **bundle exec**, but it is necessary on some systems, so I'll include it for completeness.

With that, we can run the local web server in a separate tab (Figure 1.7) as follows:[4]

```
$ rails server -b $IP -p $PORT     # Use only `rails server` if running locally
```

Now the toy application should be available on the local server as described in Section 1.3.2. (If you're using the cloud IDE, be sure to open the resulting development server in a new *browser* tab, not inside the IDE itself.)

Box 2.1. Rake

In the Unix tradition, the *make* utility has played an important role in building executable programs from source code; many a computer hacker has committed to muscle memory the line

```
$ ./configure && make && sudo make install
```

---

[4]The **rails** script is designed so that you don't need to use **bundle exec**.

commonly used to compile code on Unix systems (including Linux and Mac OS X).

Rake is *Ruby make*, a make-like language written in Ruby. Rails uses Rake extensively, especially for the innumerable little administrative tasks necessary when developing database-backed web applications. The **rake db:migrate** command is probably the most common, but there are many others; you can see a list of database tasks using **-T db**:

```
$ bundle exec rake -T db
```

To see all the Rake tasks available, run

```
$ bundle exec rake -T
```

The list is likely to be overwhelming, but don't worry, you don't have to know all (or even most) of these commands. By the end of the *Rails Tutorial*, you'll know all the most important ones.

2.2.1 A user tour

If we visit the root URL at / (read "slash", as noted in Section 1.3.4), we get the same default Rails page shown in Figure 1.9, but in generating the Users resource scaffolding we have also created a large number of pages for manipulating users. For example, the page for listing all users is at /users, and the page for making a new user is at /users/new. The rest of this section is dedicated to taking a whirlwind tour through these user pages. As we proceed, it may help to refer to Table 2.1, which shows the correspondence between pages and URLs.

We start with the page to show all the users in our application, called index; as you might expect, initially there are no users at all (Figure 2.4).

To make a new user, we visit the new page, as shown in Figure 2.5. (Since the http://0.0.0.0:3000 or cloud IDE part of the address is implicit whenever we

URL	Action	Purpose
/users	`index`	page to list all users
/users/1	`show`	page to show user with id `1`
/users/new	`new`	page to make a new user
/users/1/edit	`edit`	page to edit user with id `1`

Table 2.1: The correspondence between pages and URLs for the Users resource.

Figure 2.4: The initial index page for the Users resource (/users).

Figure 2.5: The new user page (/users/new).

are developing locally, I'll omit it from now on.) In Chapter 7, this will become the user signup page.

We can create a user by entering name and email values in the text fields and then clicking the Create User button. The result is the user `show` page, as seen in Figure 2.6. (The green welcome message is accomplished using the *flash*, which we'll learn about in Section 7.4.2.) Note that the URL is /users/1; as you might suspect, the number **1** is simply the user's **id** attribute from Figure 2.2. In Section 7.1, this page will become the user's profile.

To change a user's information, we visit the `edit` page (Figure 2.7). By modifying the user information and clicking the Update User button, we arrange to change the information for the user in the toy application (Figure 2.8).

Figure 2.6: The page to show a user (/users/1).

Figure 2.7: The user edit page (/users/1/edit).

(As we'll see in detail starting in Chapter 6, this user data is stored in a database back-end.) We'll add user edit/update functionality to the sample application in Section 9.1.

Now we'll create a second user by revisiting the `new` page and submitting a second set of user information; the resulting user `index` is shown in Figure 2.9. Section 7.1 will develop the user index into a more polished page for showing all users.

Having shown how to create, show, and edit users, we come finally to destroying them (Figure 2.10). You should verify that clicking on the link in Figure 2.10 destroys the second user, yielding an index page with only one user. (If it doesn't work, be sure that JavaScript is enabled in your browser;

Figure 2.8: A user with updated information.

Figure 2.9: The user index page (/users) with a second user.

Figure 2.10: Destroying a user.

Rails uses JavaScript to issue the request needed to destroy a user.) Section 9.4 adds user deletion to the sample app, taking care to restrict its use to a special class of administrative users.

2.2.2   MVC in action

Now that we've completed a quick overview of the Users resource, let's examine one particular part of it in the context of the Model-View-Controller (MVC) pattern introduced in Section 1.3.3. Our strategy will be to describe the results of a typical browser hit—a visit to the user index page at /users—in terms of MVC (Figure 2.11).
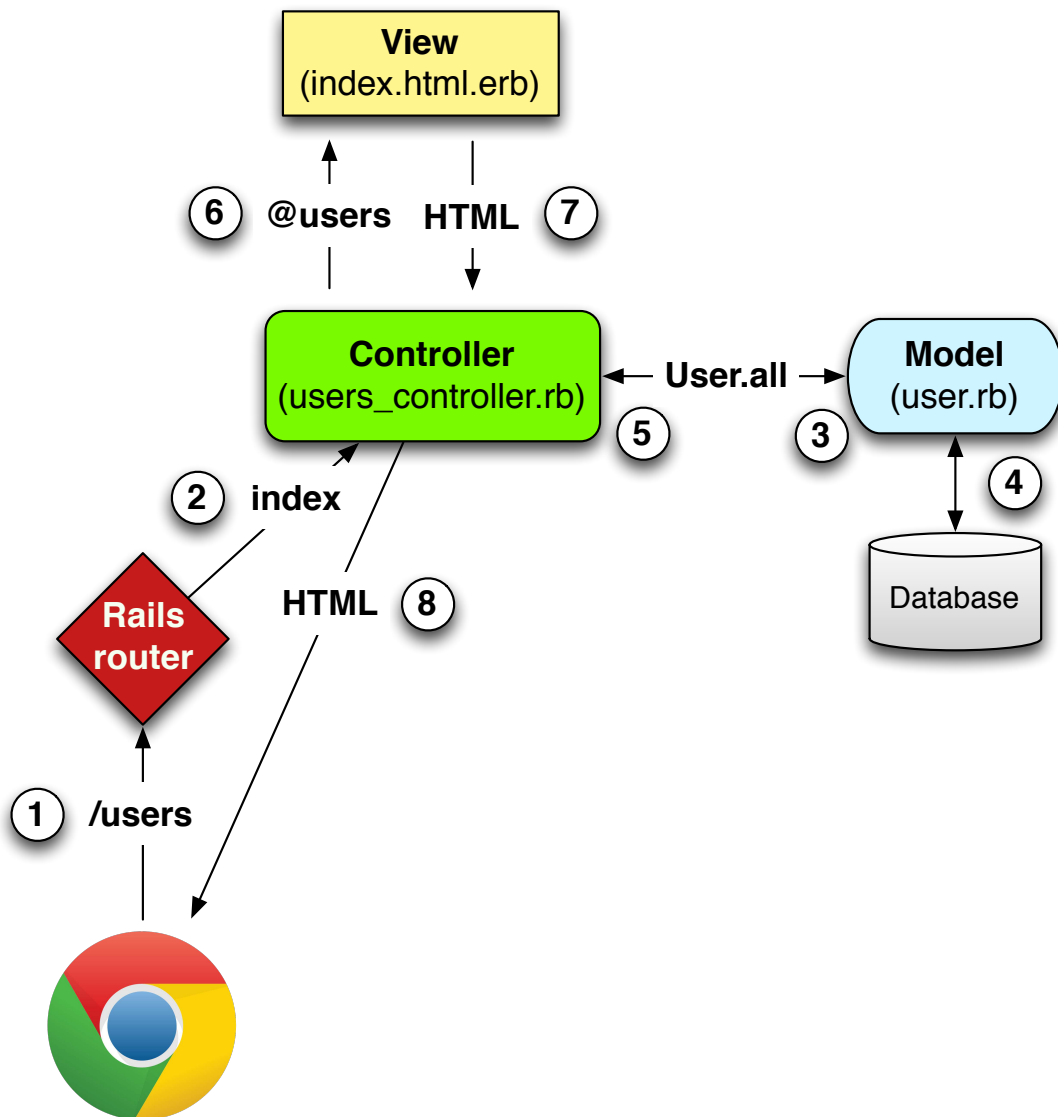
Figure 2.11: A detailed diagram of MVC in Rails.

Here is a summary of the steps shown in Figure 2.11:

1. The browser issues a request for the /users URL.

2. Rails routes /users to the **index** action in the Users controller.

3. The **index** action asks the User model to retrieve all users (**User.all**).

4. The User model pulls all the users from the database.

5. The User model returns the list of users to the controller.

6. The controller captures the users in the **@users** variable, which is passed to the **index** view.

7. The view uses embedded Ruby to render the page as HTML.

8. The controller passes the HTML back to the browser.[5]

Now let's take a look at the above steps in more detail. We start with a request issued from the browser—i.e., the result of typing a URL in the address bar or clicking on a link (Step 1 in Figure 2.11). This request hits the *Rails router* (Step 2), which dispatches to the proper *controller action* based on the URL (and, as we'll see in Box 3.2, the type of request). The code to create the mapping of user URLs to controller actions for the Users resource appears in Listing 2.2; this code effectively sets up the table of URL/action pairs seen in Table 2.1. (The strange notation **:users** is a *symbol*, which we'll learn about in Section 4.3.3.)

Listing 2.2: The Rails routes, with a rule for the Users resource.
`config/routes.rb`

---

[5]Some references indicate that the view returns the HTML directly to the browser (via a web server such as Apache or Nginx). Regardless of the implementation details, I prefer to think of the controller as a central hub through which all the application's information flows.